



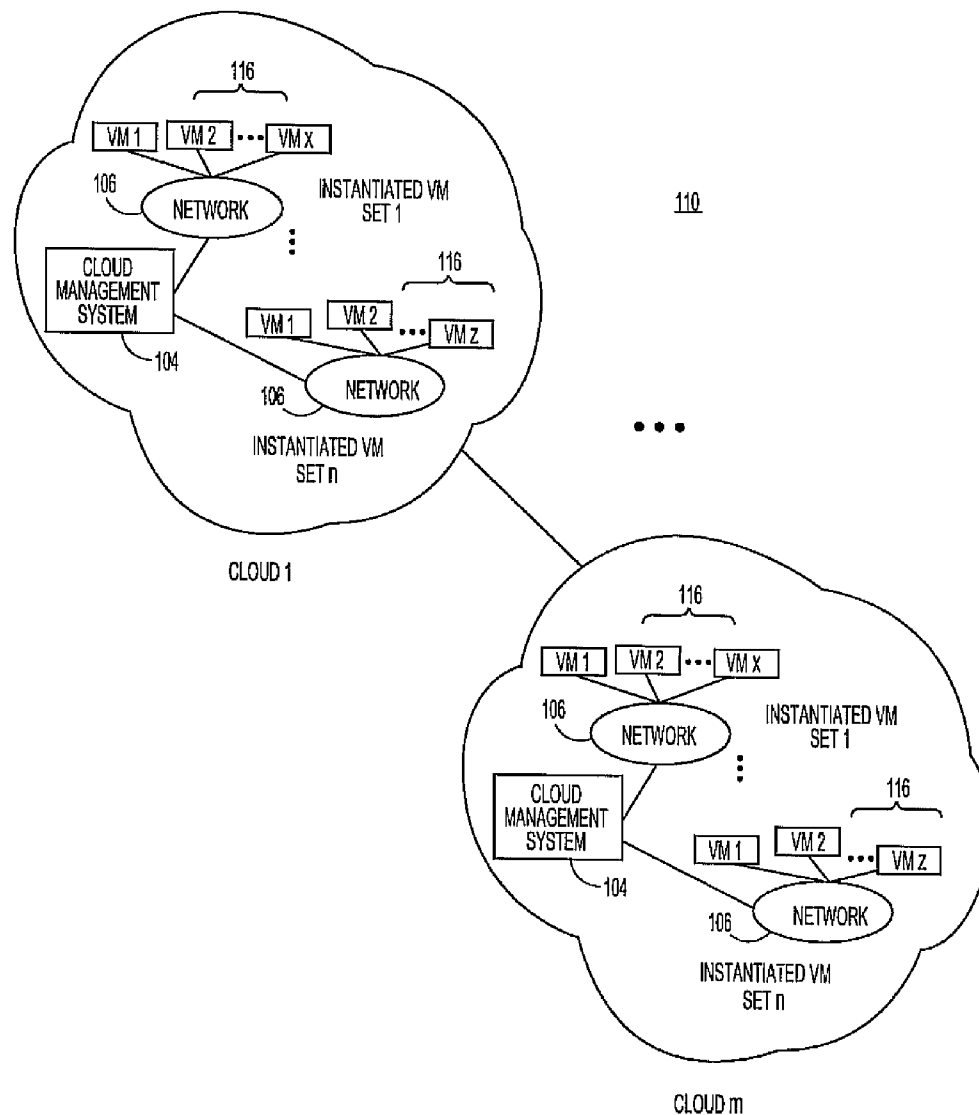
US 20100050172A1

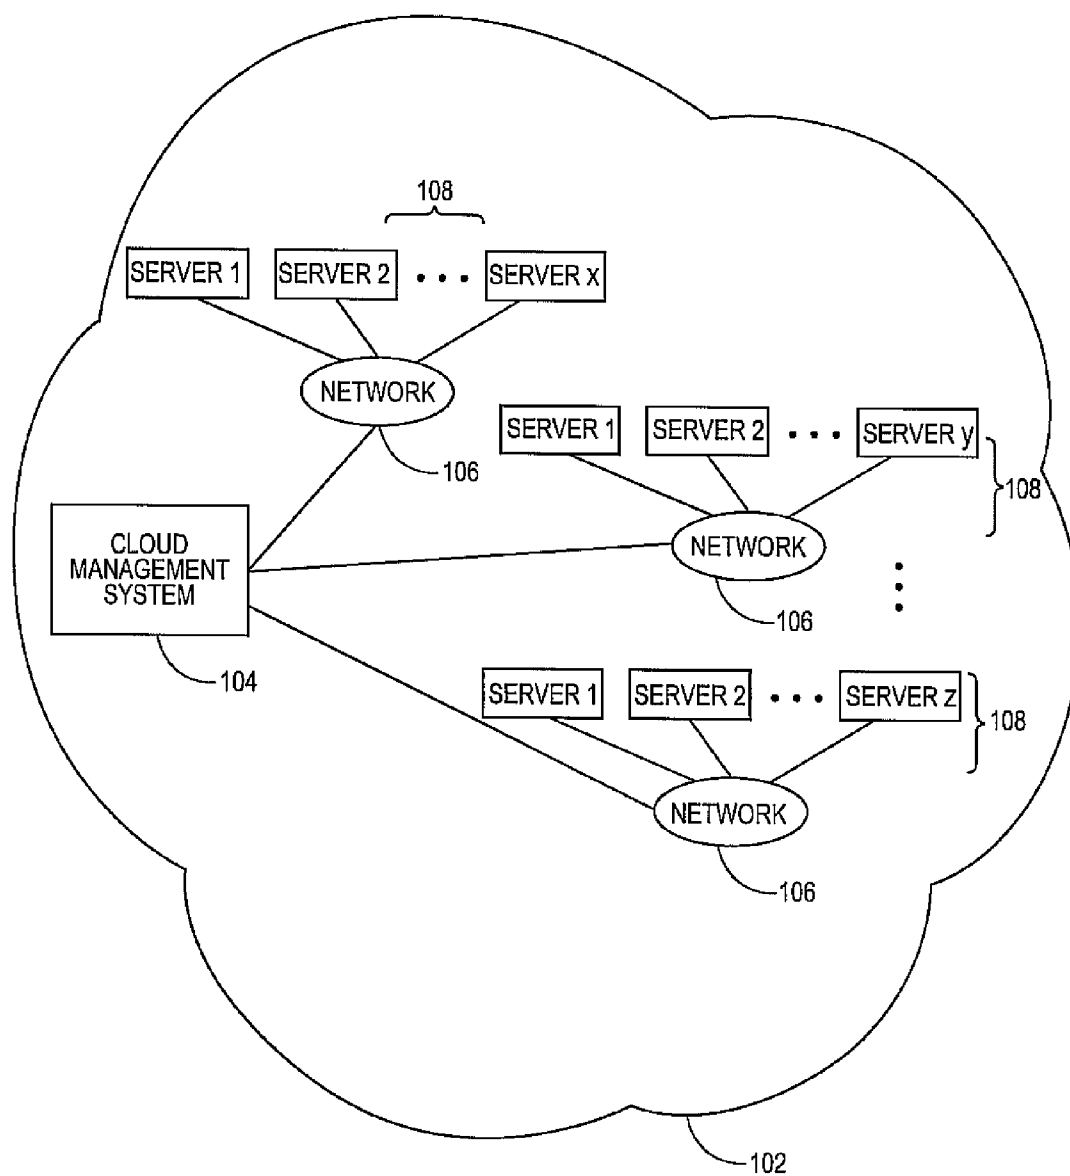
(19) **United States**(12) **Patent Application Publication**  
**Ferris**(10) **Pub. No.: US 2010/0050172 A1**(43) **Pub. Date: Feb. 25, 2010**(54) **METHODS AND SYSTEMS FOR OPTIMIZING  
RESOURCE USAGE FOR CLOUD-BASED  
NETWORKS****Publication Classification**(51) **Int. Cl.**  
**G06F 9/455** (2006.01)(52) **U.S. Cl.** ..... 718/1(57) **ABSTRACT**(76) **Inventor:** **James Michael Ferris**, Cary, NC  
(US)

Correspondence Address:

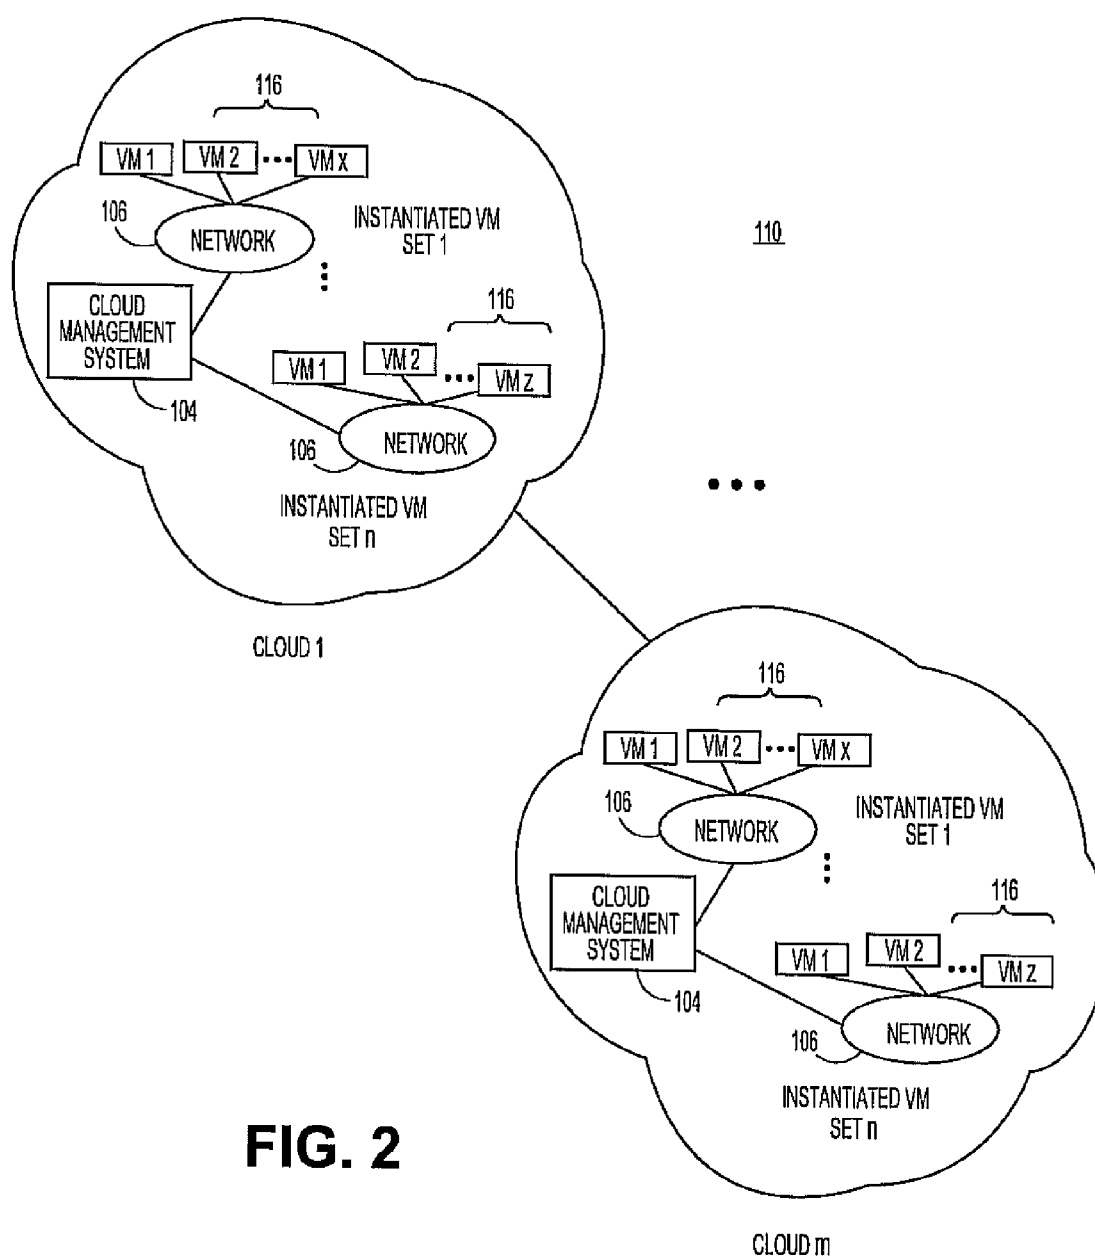
**MH2 TECHNOLOGY LAW GROUP (Cust. No.  
w/Red Hat)****1951 KIDWELL DRIVE, SUITE 550  
TYSONS CORNER, VA 22182 (US)**(21) **Appl. No.: 12/196,459**(22) **Filed: Aug. 22, 2008**

A cloud marketplace system can be configured to communicate with multiple cloud computing environments in order to ascertain the details for the resources and services provided by the cloud computing environments for optimizing resources utilized by virtual machines. The cloud marketplace system can be configured to determine the resource and service data for the cloud computing environments and select a set of resource servers for instantiating the virtual machines based specifications of the virtual machines and parameters of the instantiation. The cloud marketplace system can be configured to periodically monitor the cloud's resources and migrate the virtual machines if resources become available that more closely match the parameters of the virtual machines.





**FIG. 1**



**FIG. 2**



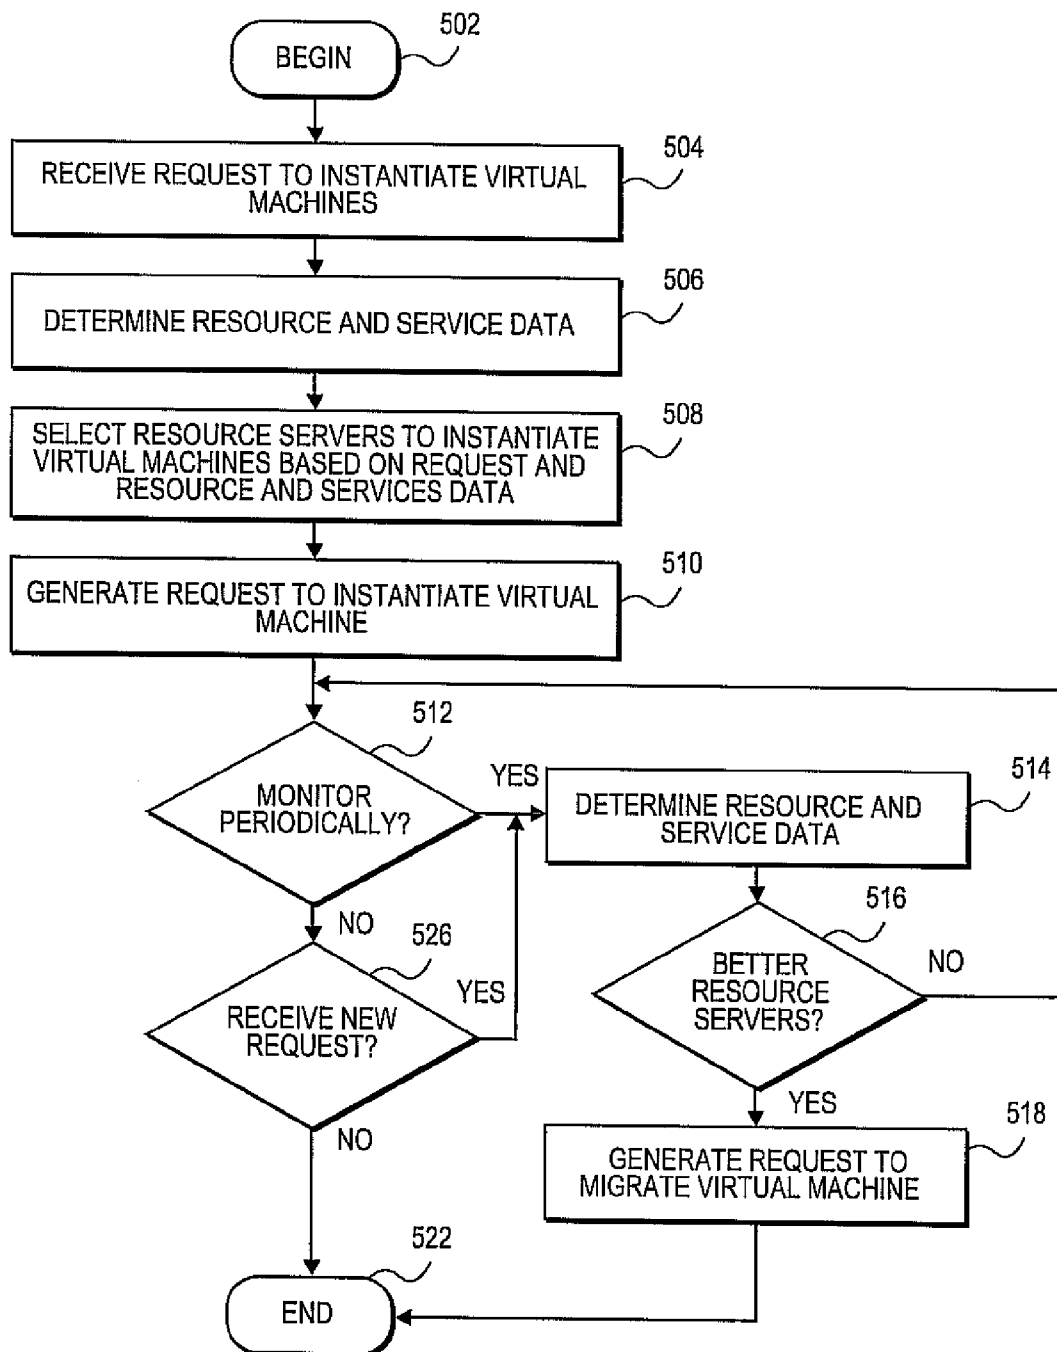


FIG. 3



**FIG. 4**



**FIG. 5**

## METHODS AND SYSTEMS FOR OPTIMIZING RESOURCE USAGE FOR CLOUD-BASED NETWORKS

### FIELD

[0001] This invention relates generally to products and services, more particularly, to systems and methods for cloud computing related services and products.

### DESCRIPTION OF THE RELATED ART

[0002] The advent of cloud-based computing architectures has opened new possibilities for the rapid and scalable deployment of virtual Web stores, media outlets, and other on-line sites or services. In general, a cloud-based architecture deploys a set of hosted resources such as processors, operating systems, software and other components that can be combined or strung together to form virtual machines. A user or customer can request the instantiation of a virtual machine or set of machines from those resources from a central server or management system to perform intended tasks or applications. For example, a user may wish to set up and instantiate a virtual server from the cloud to create a storefront to market products or services on a temporary basis, for instance, to sell tickets to an upcoming sports or musical performance. The user can lease or subscribe to the set of resources needed to build and run the set of instantiated virtual machines on a comparatively short-term basis, such as hours or days, for their intended application.

[0003] Currently, when a user desires to can lease or subscribe to the set of resources, the user may request the set of resources based on their service level requirements and then receive a price for the request. Currently, no mechanism exists for users to view different pricing information for different possible service levels. Additionally, no mechanism exist for users to view pricing information for service levels across multiple cloud computing environments. Also, no mechanism exist for users to automatically select and modify or optimize the resources utilized based on updated pricing information and updated service level information in the cloud computing environment. Thus, there is a need in the art for methods and systems that provide relevant service level and financial information and optimization services for different cloud computing environments.

### BRIEF DESCRIPTION OF THE DRAWINGS

[0004] Various features of the embodiments can be more fully appreciated, as the same become better understood with reference to the following detailed description of the embodiments when considered in connection with the accompanying figures, in which:

[0005] FIG. 1 illustrates an overall cloud system architecture in which various embodiments of the present teachings can be practiced;

[0006] FIG. 2 illustrates an overall cloud system architecture in which various embodiments of the present teachings can be practiced in another regard including multiple cloud arrangements, according to various embodiments;

[0007] FIG. 3 illustrates a network configuration in which a marketplace system can communicate with clouds, according to various embodiments;

[0008] FIG. 4 illustrates an exemplary hardware configuration for a marketplace system, according to various embodiments; and

[0009] FIG. 5 illustrates a flowchart for resource optimization in a cloud computing environment, according to various embodiments.

### DETAILED DESCRIPTION OF EMBODIMENTS

[0010] For simplicity and illustrative purposes, the principles of the present invention are described by referring mainly to exemplary embodiments thereof. However, one of ordinary skill in the art would readily recognize that the same principles are equally applicable to, and can be implemented in, all types of information and service portals, and that any such variations do not depart from the true spirit and scope of the present invention. Moreover, in the following detailed description, references are made to the accompanying figures, which illustrate specific embodiments. Electrical, mechanical, logical and structural changes may be made to the embodiments without departing from the spirit and scope of the present invention. The following detailed description is, therefore, not to be taken in a limiting sense and the scope of the present invention is defined by the appended claims and their equivalents.

[0011] Embodiments of the present teachings relate to systems and methods for optimizing resource usage for cloud computing environments. More particularly, embodiments relate to platforms and techniques in which a cloud marketplace system automatically evaluates the resources and services provided by the cloud computing environments in order to instantiate virtual machines optimized according to a requester's specifications and parameters.

[0012] According to embodiments, in general, the cloud marketplace system can be configured to communicate with multiple cloud computing environments in order to ascertain the details for the resources and services provided by the cloud computing environments for optimizing resources utilized by virtual machines. In particular, an optimization engine can be configured to receive a request to instantiate virtual machines. The request can include the specifications of the virtual machines and parameters for instantiating the virtual machines such as cost, time period, and resources/services desired.

[0013] In embodiments, the optimization engine can be configured to determine the resource and service data for the cloud computing environments. The resource and service data can include computing resources and services provided by the cloud computing environments, computing resources and services available in the cloud computing environments at particular time periods, cost data for computing resources and services provided by the cloud computing environments, and status data for the cloud computing environments at particular time periods.

[0014] In embodiments, the optimization engine can be configured to select a set of resource servers for instantiating the virtual machines based on the determined resource and service data, the specifications of the virtual machines, and the parameters of the request. Once selected, the optimization engine can request instantiation of the virtual machines or directly instantiate the virtual machines.

[0015] In embodiments, the cloud marketplace system can be configured to periodically monitor the cloud's resources and migrate the virtual machines if resources become available that more closely match the parameters of the virtual machines. In particular, the optimization engine can periodically monitor and determine the service and resource data for the cloud computing environments. The optimization engine

can compare the service and resource data to the original specifications and parameters of the virtual machines. If more closely matching resource servers become available, the optimization engine can migrate the virtual machines to the more closely matching resource servers.

**[0016]** By providing optimization for the services and resources of cloud computing environments, users can perform efficient cost-accounting for, shopping for, and migrations of existing and planned workloads to the best possible cloud computing environment. As such, the user can locate and evaluate the best priced and best quality service for their workloads and applications within minimum interaction with the clouds.

**[0017]** FIG. 1 illustrates an overall cloud computing environment, in systems and methods for the management of subscriptions of cloud-based virtual machines can operate, according to embodiments of the present teachings. Embodiments described herein can be implemented in or supported by a cloud network architecture. As used herein, a “cloud” can comprise a collection of resources that can be invoked to instantiate a virtual machine, process, or other resource for a limited or defined duration. As shown for example in FIG. 1, the collection of resources supporting a cloud **102** can comprise a set of resource servers **108** configured to deliver computing components needed to instantiate a virtual machine, process, or other resource. For example, one group of resource servers can host and serve an operating system or components thereof to deliver to and instantiate a virtual machine. Another group of resource servers can accept requests to host computing cycles or processor time, to supply a defined level of processing power for a virtual machine. A further group of resource servers can host and serve applications to load on an instantiation of a virtual machine, such as an email client, a browser application, a messaging application, or other applications or software. Other types of resource servers are possible.

**[0018]** In embodiments, the entire set of resource servers **108** or other hardware or software resources used to support the cloud **102** along with its instantiated virtual machines is managed by a cloud management system **104**. The cloud management system **104** can comprise a dedicated or centralized server and/or other software, hardware, and network tools that communicate via one or more networks **106** such as the Internet or other public or private network with all sets of resource servers to manage the cloud **102** and its operation. To instantiate a new set of virtual machines, a user can transmit an instantiation request to the cloud management system **104** for the particular type of virtual machine they wish to invoke for their intended application. A user can for instance make a request to instantiate a set of virtual machines configured for email, messaging or other applications from the cloud **102**. The request can be received and processed by the cloud management system **104**, which identifies the type of virtual machine, process, or other resource being requested. The cloud management system **104** can then identify the collection of resources necessary to instantiate that machine or resource. In embodiments, the set of instantiated virtual machines or other resources can for example comprise virtual transaction servers used to support Web storefronts, or other transaction sites.

**[0019]** In embodiments, the user’s instantiation request can specify a variety of parameters defining the operation of the set of virtual machines to be invoked. The instantiation request, for example, can specify a defined period of time for

which the instantiated machine or process is needed. The period of time can be, for example, an hour, a day, or other increment of time. In embodiments, the user’s instantiation request can specify the instantiation of a set of virtual machines or processes on a task basis, rather than for a pre-determined amount of time. For instance, a user could request resources until a software update is completed. The user’s instantiation request can specify other parameters that define the configuration and operation of the set of virtual machines or other instantiated resources. For example, the request can specify an amount of processing power or input/output (I/O) throughput the user wishes to be available to each instance of the virtual machine or other resource. In embodiments, the requesting user can for instance specify a service level agreement (SLA) acceptable for their application. Other parameters and settings can be used. One skilled in the art will realize that the user’s request can likewise include combinations of the foregoing exemplary parameters, and others.

**[0020]** When the request to instantiate a set of virtual machines or other resources has been received and the necessary resources to build that machine or resource have been identified, the cloud management system **104** can communicate with one or more set of resource servers **108** to locate resources to supply the required components. The cloud management system **104** can select providers from the diverse set of resource servers **108** to assemble the various components needed to build the requested set of virtual machines or other resources. It may be noted that in some embodiments, permanent storage such as hard disk arrays may not be included or located within the set of resource servers **108** available to the cloud management system **104**, since the set of instantiated virtual machines or other resources may be intended to operate on a purely transient or temporary basis. In embodiments, other hardware, software or other resources not strictly located or hosted in the cloud can be leveraged as needed. For example, other software services that are provided outside of the cloud **102** and hosted by third parties can be invoked by in-cloud virtual machines. For further example, other non-cloud hardware and/or storage services can be utilized as an extension to the cloud **102**, either on an on-demand or subscribed or decided basis.

**[0021]** With the resource requirements identified, the cloud management system **104** can extract and build the set of virtual machines or other resources on a dynamic or on-demand basis. For example, one set of resource servers **108** can respond to an instantiation request for a given quantity of processor cycles with an offer to deliver that computational power immediately and guaranteed for the next hour. A further set of resource servers **108** can offer to immediately supply communication bandwidth, for example on a guaranteed minimum or best-efforts basis. In other embodiments, the set of virtual machines or other resources can be built on a batch basis or at a particular future time. For example, a set of resource servers **108** can respond to a request for instantiation at a programmed time with an offer to deliver the specified quantity of processor cycles within a specific amount of time, such as the next 12 hours.

**[0022]** The cloud management system **104** can select group of servers in the set of resource servers **108** that match or best match the instantiation request for each component needed to build the virtual machine or other resource. The cloud management system **104** can then coordinate the integration of the completed group of servers from the set of resource servers **108**, to build and launch the requested set of virtual machines



or other resources. The cloud management system **104** can track the combined group of servers selected from the set of resource servers **108**, or other distributed resources that are dynamically or temporarily combined, to produce and manage the requested virtual machine population or other resources.

**[0023]** In embodiments, the cloud management system **104** can generate a resource aggregation table that identifies the various sets of resource servers that will be used to supply the components of the virtual machine or process. The sets of resource servers can be identified by unique identifiers such as, for instance, Internet Protocol (IP) addresses or other addresses. The cloud management system **104** can register the finalized group of servers in the set of resource servers **108** contributing to an instantiated machine or process.

**[0024]** The cloud management system **104** can then set up and launch the initiation process for the virtual machines, processes, or other resources to be delivered from the cloud. The cloud management system **104** can for instance transmit an instantiation command or instruction to the registered group of servers in set of resource servers **108**. The cloud management system **104** can receive a confirmation message back from each participating server in set of resource servers **118** indicating a status regarding the provisioning of their respective resources. Various sets of resource servers can confirm, for example, the availability of a dedicated amount of processor cycles, amounts of electronic memory, communications bandwidth, or applications or other software prepared to be served.

**[0025]** As shown for example in FIG. 2, the cloud management system **104** can then instantiate one or more than one set of virtual machines **116**, or other processes based on the resources supplied by the registered set of resource servers **108**. In embodiments, the cloud management system **104** can instantiate a given number, for example, 10, 500, 1000, or other numbers of virtual machines to be made available to users on a network **114**, such as the Internet or other public or private network. Each virtual machine can be assigned an instantiated machine ID that can be stored in the resource aggregation table, or other record or image of the instantiated population. Additionally, the cloud management system **104** can store the duration of each virtual machine and the collection of resources utilized by the complete set of instantiated virtual machines **116**.

**[0026]** In embodiments, the cloud management system **104** can further store, track and manage a user's identity and associated set of rights or entitlements to software, hardware, and other resources. Each user that populates a set of virtual machines in the cloud can have specific rights and resources assigned and made available to them. The cloud management system **104** can track and configure specific actions that a user can perform, such as provision a set of virtual machines with software applications or other resources, configure a set of virtual machines to desired specifications, submit jobs to the set of virtual machines or other host, manage other users of the set of instantiated virtual machines **116** or other resources, and other privileges or actions. The cloud management system **104** can further generate records of the usage of instantiated virtual machines to permit tracking, billing, and auditing of the services consumed by the user. In embodiments, the cloud management system **104** can for example meter the usage and/or duration of the set of instantiated virtual

machines **116**, to generate subscription billing records for a user that has launched those machines. Other billing or value arrangements are possible.

**[0027]** The cloud management system **104** can configure each virtual machine to be made available to users of the one or more networks **106** via a browser interface, or other interface or mechanism. Each instantiated virtual machine can communicate with the cloud management system **104** and the underlying registered set of resource servers **108** via a standard Web application programming interface (API), or via other calls or interfaces. The set of instantiated virtual machines **116** can likewise communicate with each other, as well as other sites, servers, locations, and resources available via the Internet or other public or private networks, whether within a given cloud **102** or between clouds.

**[0028]** It may be noted that while a browser interface or other front-end can be used to view and operate the set of instantiated virtual machines **116** from a client or terminal, the processing, memory, communications, storage, and other hardware as well as software resources required to be combined to build the virtual machines or other resources are all hosted remotely in the cloud **102**. In embodiments, the set of virtual machines **116** or other resources may not depend on or require the user's own on-premise hardware or other resources. In embodiments, a user can therefore request and instantiate a set of virtual machines or other resources on a purely off-premise basis, for instance to build and launch a virtual storefront or other application.

**[0029]** Because the cloud management system **104** in one regard specifies, builds, operates and manages the set of instantiated virtual machines **116** on a logical level, the user can request and receive different sets of virtual machines and other resources on a real-time or near real-time basis, without a need to specify or install any particular hardware. The user's set of instantiated virtual machines **116**, processes, or other resources can be scaled up or down immediately or virtually immediately on an on-demand basis, if desired. In embodiments, the various sets of resource servers that are accessed by the cloud management system **104** to support a set of instantiated virtual machines **116** or processes can change or be substituted, over time. The type and operating characteristics of the set of instantiated virtual machines **116** can nevertheless remain constant or virtually constant, since instances are assembled from abstracted resources that can be selected and maintained from diverse sources based on uniform specifications.

**[0030]** In terms of network management of the set of instantiated virtual machines **116** that have been successfully configured and instantiated, the cloud management system **104** can perform various network management tasks including security, maintenance, and metering for billing or subscription purposes. The cloud management system **104** of a given cloud **102** can, for example, install or terminate applications or appliances on individual machines. The cloud management system **104** can monitor operating virtual machines to detect any virus or other rogue process on individual machines, and for instance terminate the infected application or virtual machine. The cloud management system **104** can likewise manage an entire set of instantiated virtual machines **116** or other resources on a collective basis, for instance, to push or delivery a software upgrade to all active virtual machines. Other management processes are possible.

**[0031]** In embodiments, more than one set of virtual machines can be instantiated in a given cloud at the same,

overlapping or successive times. The cloud management system **104** can, in such implementations, build, launch and manage multiple sets of virtual machines based on the same or different underlying set of resource servers **108**, with populations of different sets of instantiated virtual machines **116** such as may be requested by different users. The cloud management system **104** can institute and enforce security protocols in a cloud **102** hosting multiple sets of virtual machines. Each of the individual sets of virtual machines can be hosted in a respective partition or sub-cloud of the resources of the cloud **102**. The cloud management system **104** of a cloud can for example deploy services specific to isolated or defined sub-clouds, or isolate individual workloads/processes within the cloud to a specific sub-cloud. The subdivision of the cloud **102** into distinct transient sub-clouds or other sub-components which have assured security and isolation features can assist in establishing a multiple user or multi-tenant cloud arrangement. In a multiple user scenario, each of the multiple users can use the cloud platform as a common utility while retaining the assurance that their information is secure from other users of the overall cloud system. In further embodiments, sub-clouds can nevertheless be configured to share resources, if desired.

[0032] In embodiments, and as also shown in FIG. 2, the set of instantiated virtual machines **116** generated in a first cloud **102** can also interact with a set of instantiated virtual machines or processes generated in a second, third or further cloud **102**. The cloud management system **104** of a first cloud **102** can interface with the cloud management system **104** of a second cloud **102**, to coordinate those domains and operate the clouds and/or virtual machines or processes on a combined basis. The cloud management system **104** of a given cloud **102** can track and manage individual virtual machines or other resources instantiated in that cloud, as well as the set of instantiated virtual machines or other resources in other clouds.

[0033] In the foregoing and other embodiments, the user making an instantiation request or otherwise accessing or utilizing the cloud network can be a person, customer, subscriber, administrator, corporation, organization, or other entity. In embodiments, the user can be or include another virtual machine, application or process. In further embodiments, multiple users or entities can share the use of a set of virtual machines or other resources.

[0034] FIG. 3 illustrates aspects of a cloud marketplace system **200** and its interaction with various clouds **102**, according to various embodiments. In embodiments, the cloud marketplace system **200** provides a marketplace for the services provided by clouds **102**. In embodiments as shown, the cloud marketplace system **200** can comprise a marketplace engine **205**, an optimization engine **210**, a marketplace store **215**, and an application programming interface (API) **220**. The marketplace engine **205** can be configured to contain control logic for performing the marketplace processes and for communicating with the clouds **102**, API **220**, and marketplace store **215**. In embodiments as shown, the cloud marketplace system **200** provides a marketplace for the services provided by clouds **102**.

[0035] In embodiments, the cloud marketplace system **200** can be configured to communicate with the clouds **102** in order to ascertain the details for the resources and services provided by the clouds **102**. In particular, the cloud marketplace system **200** can be configured to retrieve, from the clouds **102**, resource and service data **225** which describes the

capabilities, services, status, and cost of the clouds **102**. The resource and service data **225** can include data detailing computing resources and services provided by the clouds **102**, computing resources and services available in the clouds **102** at particular time periods, cost for the services provided by the clouds **102**, and status for the clouds **102** at particular time periods.

[0036] In embodiments, the data detailing the computing resources and services provided by the clouds **102** can include resource data specifying the sets of resources servers **108** contained in the clouds **102** and the computing resources provided by the sets of resources servers **108** (computing cycles, bandwidth, memory, storage). Additionally, the data detailing the computing resources and services provided by the clouds **102** can include service data specifying particular services provided by the clouds **102** such as identity management services, cloud management services, application and appliance building services, and the like.

[0037] In embodiments, the data detailing the computing resources and services available in the clouds **102** at particular time periods can include data specifying computing resources and services available in the clouds **102** current or at some future time. Additionally, the data detailing the computing resources and services available in the clouds **102** at particular time periods can include data specifying service level agreements for the cloud **102** for a period of time.

[0038] In embodiments, the data detailing the cost for the resources and services provided by the clouds **102** can include data such as the cost for subscription to the clouds **102**, the cost for usage of the resources and services of the clouds **102**, and any discounts of the costs based on the usage. For example, the cost for usage can include the cost for the resources consumed to support the virtual machines and can include the cost for the utilization of the virtual machines by the user or third parties. These can include the cost for the instantiated time of the virtual machines, the computing resource consumed by the virtual machines (computing cycles, bandwidth), the time frame of the computing resource consumption (peak time, off-peak time), and combinations thereof. The cost data can include cost data for a current time or cost data for resources and services in the future, such as total lifetime costs of a cloud application.

[0039] In embodiments, the data detailing the status for the clouds **102** at particular time periods can include data specifying the resources and services current being utilized in the clouds **102**. For example, the status data can include virtual machines currently instantiated in the clouds **102** and the computing resources of the sets of resources servers **108** (computing cycles, bandwidth, memory, storage) currently utilized.

[0040] In embodiments, in order to acquire the resource and service data **225**, the cloud marketplace system **200** can be configured to request the resource and service data **225** from cloud management systems **104** located in the clouds **102**. Likewise, the cloud marketplace system **200** can be configured to communicate with the set of resource servers **108** directly in order to obtain the resource and service data **225**. The cloud marketplace system **200** can be configured to communicate with the clouds **102** via network **106** or any other available public or private network. In particular, marketplace engine **205** can be configured to communicate with the clouds **102**.

[0041] In embodiments, the cloud marketplace system **200** can be configured to maintain a repository **230** for the

resource and service data **225** in order to provide the marketplace for the clouds **102**. In particular, the marketplace engine **205** can be configured to maintain the repository **230**. The repository **230** can be configured in any format to allow the resource and service data **225** to be stored and retrieved. For example, the repository can be configured in any type of proprietary or open-source database format. The repository **230** can include a record for each cloud **102** that includes various searchable data fields for storing the resource and service data **225** for that particular cloud.

[0042] The cloud marketplace system **200** can be configured to update the repository **230** periodically to maintain current resource and service data **225** for the clouds **102**. Likewise, when providing the marketplace, the cloud marketplace system **200** can be configured to obtain the resource and service data **225**, on-demand, and directly from the clouds **102**.

[0043] In embodiments, to provide the marketplace, the cloud marketplace system **200** can be configured receive a request for information pertaining to the resources or services provided by or available in the clouds **102**. For example, a initiator may request information on clouds **102** that can provide a certain level of computing resources at a specific cost during a specific time period. The initiator of the request the information in order to see all available resources of the clouds **102** that may meet the initiator's requirements. As such, the request can include the information regarding the resources desired by the initiator. This can include the desired computing resources, the time period for the resources, the duration of the resources, the cost of the resources, services provided by the cloud, the availability of the resources, and the like.

[0044] In particular, the marketplace engine **205** can be configured to communicate with the initiator of the request via API **220**. API **220** can be configured to generate GUIs, e.g. dialog boxes, web pages, as required by marketplace engine **205** and to provide an interface for receiving requests. The initiator of the request can be a user desiring utilization of the clouds **102** or the cloud management systems **104** of the clouds **102**.

[0045] Once the request is received, the cloud marketplace system **200** can be configured to search the repository **230** for resource and service data **225** relating to the request. In particular, the marketplace engine **205** can be configured to access the repository **230** stored in marketplace store **215** and search the repository **230** for resource and service data **225** matching the request. Likewise, the cloud marketplace system **200** can be configured to obtain the resource and service data **225** directly from the clouds **102**.

[0046] Once located, the cloud marketplace system **200** can be configured to generate a marketplace report **235** detailing the resource and service data **225** matching the request. For example, the marketplace report **235** can include a list of clouds **102** matching the request along with the data detailing computing resources and services provided by the matching clouds **102**, computing resources and services available in the matching clouds **102** at particular time periods specified in the request, cost for the services and resources requested, and status for the matching clouds **102** at particular time periods specified in the request.

[0047] In particular, the marketplace engine **205** can be configured to extract the matching resource and service data

from repository **230**. Marketplace engine **205** can be configured to arrange the matching resource and service data **225** in marketplace report **235**.

[0048] After generating the marketplace report **235**, the cloud marketplace system **200** can be configured to provide the marketplace report **235** to the initiator of the request. In particular, the marketplace engine **205** can be configured to provide the marketplace report **235** via API **220**.

[0049] Once the report is received, the initiator or the user can make a determination on which resources of the clouds **102** to utilize. This can include utilizing resources from a single cloud **102** or resources from multiple clouds **102**. For example, the initiator can select one or more cloud **102** with which to instantiate a virtual machine or machines. The initiator can make a request for the resources directly to the clouds **102**. Likewise, the cloud marketplace system **200** can be configured to receive the initiators requests for resources and forward the request to the clouds **102**, or reserve the resources directly.

[0050] In additional embodiments, the cloud marketplace system **200** can be configured to utilize the resource and service data **225** to provide automatic and optimized initiation and migration services for virtual machines initiated in the clouds **102**. In particular, the optimization engine **210** can be configured to contain control logic for performing the initiation and migration processes and for communicating with the clouds **102**, the marketplace engine **205**, API **220**, and the marketplace store **215**.

[0051] In embodiments, for a particular virtual machine or machines to be instantiated in the cloud **102**, the cloud marketplace system **200** can be configured to receive a request to instantiate virtual machines in the clouds **102**. The request can include specifications of the virtual machines or copies of the virtual machines. The request can also include parameters for instantiating the virtual machines, such as desired cost for instantiating the virtual machine, time period for instantiation of the virtual machine, desired resources for the virtual machine (computing cycles, bandwidth, etc), and the like.

[0052] In embodiments, once the request is received, the optimization engine **210** can be configured to determine a set of resource servers from clouds **102** on which to instantiate the virtual machines according to the request. In particular, the optimization engine **210** can be configured to determine a set of resource servers that most closely match the specifications for the virtual machines and the parameters specified by the requester. For example, a requester can request that virtual machines be instantiated in order to minimize cost and, in response, the optimization engine **210** can determine a set of resources servers that meet the specifications of the virtual machines while minimizing the cost of instantiating and supporting the virtual machines.

[0053] In embodiments, in order to determine a set of resources servers, the optimization engine **210** can be configured to determine resource and service data **225** for cloud **102**. The optimization engine **210** can be configured to request resource and service data **225** from marketplace engine **205**. Additionally, the optimization engine **210** can be configured to retrieve stored resource and service data **225** from marketplace store **215**. Also, the optimization engine **210** can be configured to include the logic of the marketplace engine **205** in order to determine resource and service data **225** directly from the clouds **102**. As mentioned above, the resource and service data **225** can describe the capabilities, services, status, and cost of the clouds **102**. The resource and service data **225**

can include data detailing computing resources and services provided by the clouds **102**, computing resources and services available in the clouds **102** at particular time periods, cost for the services provided by the clouds **102**, and status for the clouds **102** at particular time periods.

**[0054]** Once the resource and service data **225** is determined, the optimization engine **210** can be configured to select a set of resource servers from clouds **102** that most closely match the specifications and parameters of the request. The optimization engine **210** can compare the specifications and parameters to resources and services data **225** in order to select a set of resource servers that most closely match the request. The optimization engine **210** can select the set of resource servers for the overall virtual machines. Likewise, the optimization engine **210** can select the set of resource servers that most closely match the request based on each component of the virtual machines. That is, the optimization engine **210** can compare the specifications and parameters for each component of the virtual machines to resources and services data **225** in order to select a set of resource servers for each component.

**[0055]** In embodiments, once the set of resource servers are selected, optimization engine **210** can be configured to send a request to the clouds **102** to initiate the virtual machines. Additionally, the optimization engine **210** can be configured to directly instantiate the virtual machines in the clouds **102**.

**[0056]** In embodiments, for a particular virtual machine or machines instantiated in the clouds **102**, the cloud marketplace system **200** can be configured to periodically monitor the clouds **102** for resources that more closely match the specification and parameters of virtual machines and migrate the virtual machines if more closely matched resources become available. In particular, the optimization engine **210** can be configured to periodically monitor resource and service data **225** and compare resources and service data **225** to the specification and parameters of the virtual machines. If a more closely matched set of resource servers become available, the optimization engine **210** can be configured to migrate the virtual machines to the more closely matched set of resource servers in clouds **102**. The optimization engine **210** can be configured to migrate the entire virtual machines or portions of the virtual machines.

**[0057]** To migrate the virtual machine or machines, the optimization engine **210** can be configured send a request, to the clouds **102** currently instantiating the virtual machine or machines, to migrate the virtual machine or machines to a new set of resource servers. The request can include an identification of the virtual machine or machines and an identification of the new set of resource servers. Likewise, the optimization engine **210** can be configured to retrieve the virtual machine or machines from the clouds **102** currently instantiating the virtual machine or machines and pass the virtual machine or machines to the new set of resource servers.

**[0058]** Prior to migrating, the cloud marketplace system **200** can be configured to notify the user or initiator of the possible migration and allow the initiator to approve or decline the migration. Likewise, the cloud marketplace system **200** can be configured to migrate the virtual machines automatically and transparent to the user or initiator. Once migrated, cloud marketplace system **200** can be configured to notify the user or initiator of the migration. For notification, the optimization engine **210** can be configured to notify the user or initiator via API **220**.

**[0059]** In embodiments, once virtual machines are instantiated, the cloud marketplace system **200** can be configured to receive updated parameters for the instantiated virtual machines. For example, a requester can send a request to

change the virtual machines from being optimized for cost to being optimized for a particular resource or service level. As such, the optimization engine **210** can be configured to determine the resource and service data **225**, as described above, and compare resources and service data **225** to the new parameters of the virtual machines. If a more closely matched set of resource servers are available, the optimization engine **210** can be configured to migrate the virtual machines to the more closely matched set of resource servers in the clouds **102**, as described above.

**[0060]** FIG. 4 illustrates an exemplary diagram of hardware and other resources that can be incorporated in a cloud marketplace system **200** configured to communicate with instantiated the clouds **102** and users or initiators via one or more networks **106**, according to embodiments. In embodiments as shown, the cloud marketplace system **200** can comprise a processor **300** communicating with memory **302**, such as electronic random access memory, operating under control of or in conjunction with operating system **306**. Operating system **306** can be, for example, a distribution of the Linux™ operating system, the Unix™ operating system, or other open-source or proprietary operating system or platform. Processor **300** also communicates with the marketplace store **215**, such as a database stored on a local hard drive. Processor **300** further communicates with network interface **304**, such as an Ethernet or wireless data connection, which in turn communicates with one or more networks **106**, such as the Internet or other public or private networks. Processor **300** also communicates with marketplace store **215** and the marketplace engine **205**, to execute control logic and perform the marketplace process described above. Additionally, processor **300** also communicates with marketplace store **215** and the optimization engine **210**, to execute control logic and perform the migration process described above. Other configurations of the cloud marketplace system **200**, associated network connections, and other hardware and software resources are possible.

**[0061]** While FIG. 4 illustrates the cloud marketplace system **200** as a standalone system comprises a combination of hardware and software, the cloud marketplace system **200** can also be implemented as a software application or program capable of being executed by a convention computer platform. Likewise, the cloud marketplace system **200** can also be implemented as a software module or program module capable of being incorporated in other software applications and programs. In either case, the cloud marketplace system **200** can be implemented in any type of conventional proprietary or open-source computer language.

**[0062]** Additionally, as illustrated in FIG. 4, the cloud marketplace system **200** can be a standalone system capable of being access by or controlled by a user or initiator of the marketplace, optimization, and migration processes. Likewise, the cloud marketplace system **200** can be integrated into one or more of the clouds **102**. Additionally, the cloud marketplace system **200** can be integrated into one or more of the systems of the clouds **102** such as the cloud management systems.

**[0063]** FIG. 5 illustrates a flow diagram of overall optimization and migration processes in a cloud marketplace system **200**, according to embodiments of the present teachings. In **502**, processing can begin. In **504**, the optimization engine **210** can receive a request to instantiate virtual machines in the clouds **102**. In **506**, the optimization engine **210** can determine the resource and service data **225** which describes the capabilities, services, status, and cost of the clouds **102**. The resource and service data can include computing resources and services provided by the clouds **102**, computing

resources and services available in the clouds **102** at particular time periods, cost data for computing resources and services provided by the clouds **102**, and status data for the clouds **102** at particular time periods.

[0064] In order to acquire the resource and service data, the optimization engine **210** can request the resource and service data from marketplace engine **205** or marketplace store **215**. Likewise, the optimization engine **210** can communicate with the resources of the clouds **102** directly in order to obtain the resource and service data.

[0065] In **508**, the optimization engine **210** can select a set of resource servers that match the request based on the resource and service data **225**. In **510**, the optimization engine **210** can generate a request to instantiate the virtual machines on the selected set of resource servers.

[0066] In **512**, the optimization engine **210** can periodically monitor resource and services data **225** for the clouds **102**. In **514**, the optimization engine **210** can periodically determine the resource and service data **225**. In **516**, the optimization engine **210** can determine if an additional set of resource servers exists that more closely match the original request. If an additional set of resource servers exists, in **518**, the optimization engine **210** can migrate the virtual machines to the additional set of resource servers. Otherwise, the optimization engine **210** can continue to monitor the resource and service data **225** for better matching resource servers.

[0067] In **520**, optimization engine **210** can receive a request to optimize the virtual machines based on a new parameter. If the request is received, the optimization engine **210** can determine if an additional set of resource servers exists that match the new request. If an additional set of resource servers exists, the optimization engine **210** can migrate the virtual machines to the additional set of resource servers. Then, in **522**, the process can end, but the process can return to any point and repeat.

[0068] In the foregoing and other embodiments, the user or initiator making a request for marketplace processes or migration processes can be a person, customer, subscriber, corporation, organization, or other entity. In embodiments, the user can be or include another virtual machine, application or process. In further embodiments, multiple users or entities can share the use of a set of virtual machines or other resources.

[0069] While the invention has been described with reference to the exemplary embodiments thereof, those skilled in the art will be able to make various modifications to the described embodiments without departing from the true spirit and scope. The terms and descriptions used herein are set forth by way of illustration only and are not meant as limitations. In particular, although the method has been described by examples, the steps of the method may be performed in a different order than illustrated or simultaneously. Those skilled in the art will recognize that these and other variations are possible within the spirit and scope as defined in the following claims and their equivalents.

What is claimed is:

1. A method of providing optimized resources for cloud computing environments, comprising:

receiving a request to instantiate at least one virtual machine in a set of cloud computing environments, the request including specifications of the at least one virtual machine and at least one parameter for instantiating the virtual machine;

determining resource and service data for the set of cloud computing environments;

selecting resource servers from the set of cloud computing environments to instantiate the at least one virtual

machine based on the resource and service data, the specification of the at least one virtual machine, and the at least one parameter; and

generating a request to instantiate the at least one virtual machine on the selected resource servers.

2. The method of claim 1, wherein determining the resource and service data comprises:

obtaining the resource and service data from at least one of cloud management systems for the set of cloud computing environments or a set of resource servers in the set of cloud computing environments.

3. The method of claim 2, wherein the resource and service data is obtained from a repository.

4. The method of claim 1, wherein receiving the request comprises:

receiving the request from at least one of an initiator of the request or a cloud management system of the set of cloud computing environments.

5. The method of claim 1, wherein the resource and service data comprises at least one of computing resources and services provided by the set of cloud computing environments, computing resources and services available in the set of cloud computing environments at particular time periods, cost data for computing resources and services provided by the set of cloud computing environments, and status data for the set of cloud computing environments at particular time periods.

6. The method of claims 1, wherein the at least one parameters comprises a cost of resources in the cloud computing environments, a service level of resources in the cloud computing environment, and time period for instantiation of the at least one virtual machine.

7. The method of claim 1, further comprising:

monitoring the resource and service data periodically for additional resources servers that match the specifications of the at least one virtual machine and the at least one parameter for instantiating the virtual machine; and migrating the at least one virtual machine to the additional resource servers when the additional resource servers match the specifications of the at least one virtual machine and the at least one parameter for instantiating the virtual machine.

8. The method of claim 1, further comprising:

receiving additional parameters for the at least one virtual machine;

selecting additional resource servers from the set of cloud computing environments to migrate the at least one virtual machine based on the resource and service data, the specification of the at least one virtual machine, and the additional parameters.

9. The method of claim 1, further comprising:

instantiating the at least one virtual machine on the selected resource servers.

10. A system for providing a marketplace for cloud computing environments, comprising:

a network interface to a set of cloud computing environments; and

a cloud marketplace module, communicating with the network interface, the cloud marketplace module being configured to—

receive a request to instantiate at least one virtual machine in a set of cloud computing environments, the request including specifications of the at least one virtual machine and at least one parameter for instantiating the virtual machine;

determine resource and service data for the set of cloud computing environments;  
 select resource servers from the set of cloud computing environments to instantiate the at least one virtual machine based on the resource and service data, the specification of the at least one virtual machine, and the at least one parameter; and  
 generate a request to instantiate the at least one virtual machine on the selected resource servers.

**11.** The system of claim **10**, wherein determining the resource and service data comprises:  
 obtaining the resource and service data from at least one of cloud management systems for the set of cloud computing environments or a set of resource servers in the set of cloud computing environments.

**12.** The system of claim **11**, wherein the resource and service data is obtained from a repository.

**13.** The system of claim **10**, wherein receiving the request comprises:  
 receiving the request from at least one of an initiator of the request or a cloud management system of the set of cloud computing environments.

**14.** The system of claim **10**, wherein the resource and service data comprises at least one of computing resources and services provided by the set of cloud computing environments, computing resources and services available in the set of cloud computing environments at particular time periods, cost data for computing resources and services provided by the set of cloud computing environments, and status data for the set of cloud computing environments at particular time periods.

**15.** The system of claims **10**, wherein the at least one parameters comprises a cost of resources in the cloud computing environments, a service level of resources in the cloud computing environment, and time period for instantiation of the at least one virtual machine.

**16.** The system of claim **10**, the cloud marketplace module being further configured to—  
 monitor the resource and service data periodically for additional resources servers that match the specifications of the at least one virtual machine and the at least one parameter for instantiating the virtual machine; and  
 migrate the at least one virtual machine to the additional resource servers when the additional resource servers match the specifications of the at least one virtual machine and the at least one parameter for instantiating the virtual machine.

**17.** The system of claim **10**, the cloud marketplace module being further configured to—

receive additional parameters for the at least one virtual machine;  
 select additional resource servers from the set of cloud computing environments to migrate the at least one virtual machine based on the resource and service data, the specification of the at least one virtual machine, and the additional parameters.

**18.** The system of claim **10**, the cloud marketplace module being further configured to—  
 instantiate the at least one virtual machine on the selected resource servers.

**19.** A cloud marketplace application, the cloud marketplace application being embodied in a computer readable medium and comprising instructions for causing a computer to perform a method comprising:  
 receiving a request to instantiate at least one virtual machine in a set of cloud computing environments, the request including specifications of the at least one virtual machine and at least one parameter for instantiating the virtual machine;  
 determining resource and service data for the set of cloud computing environments;  
 selecting resource servers from the set of cloud computing environments to instantiate the at least one virtual machine based on the resource and service data, the specification of the at least one virtual machine, and the at least one parameter; and  
 generating a request to instantiate the at least one virtual machine on the selected resource servers.

**20.** The cloud marketplace application of claim **19**, wherein the method further comprises:  
 monitoring the resource and service data periodically for additional resources servers that match the specifications of the at least one virtual machine and the at least one parameter for instantiating the virtual machine; and  
 migrating the at least one virtual machine to the additional resource servers when the additional resource servers match the specifications of the at least one virtual machine and the at least one parameter for instantiating the virtual machine.

**21.** The cloud marketplace application of claim **19**, wherein the method further comprises:  
 receiving additional parameters for the at least one virtual machine;  
 selecting additional resource servers from the set of cloud computing environments to migrate the at least one virtual machine based on the resource and service data, the specification of the at least one virtual machine, and the additional parameters.

\* \* \* \* \*