

2014 IEEE 26th International Symposium on Computer Architecture and High Performance Computing

(SBAC-PAD 2014)

**Paris, France
22 - 24 October 2014**



**IEEE Catalog Number: CFP14307-POD
ISBN: 978-1-4799-6906-7**

2014 IEEE 26th International Symposium on Computer Architecture and High Performance Computing SBAC-PAD 2014

Table of Contents

Message from the General Chairs.....	x
Message from the Program and Track Chairs.....	xi
Committee Lists.....	xiv
Keynotes.....	xviii

Shared-Memory and Multi-core Systems

A Case Study of Hybrid Dataflow and Shared-Memory Programming Models: Dependency-Based Parallel Game Engine	1
<i>Vladimir Gajinov, Igor Erić, Saša Stojanović, Veljko Milutinović, Osman Unsal, Eduard Ayguadé, and Adrián Cristal</i>	
Wait-Free Global Virtual Time Computation in Shared Memory TimeWarp Systems	9
<i>Alessandro Pellegrini and Francesco Quaglia</i>	
Compact Hash Tables for High-Performance Traffic Classification on Multi-core Processors	17
<i>Yun R. Qu and Viktor K. Prasanna</i>	

Performance Optimization

Flying Memcache: Lessons Learned from Different Acceleration Strategies	25
<i>Dimitris Deyannis, Lazaros Koromilas, Giorgos Vasiliadis, Elias Athanasopoulos, and Sotiris Ioannidis</i>	
Improving an MPI Application-Level Migration Approach through Checkpoint File Splitting	33
<i>Mónica Rodríguez, Iván Cores, Patricia González, and María J. Martín</i>	
HPCG: Preliminary Evaluation and Optimization on Tianhe-2 CPU-only Nodes	41
<i>Cheng Chen, Yunfei Du, Hao Jiang, Ke Zuo, and Canqun Yang</i>	

Leveraging Optimization Methods for Dynamically Assisted Control-Flow Integrity Mechanisms	49
<i>João Moreira, Lucas Teixeira, Edson Borin, and Sandro Rigo</i>	

Energy Aware Computing

Energy Efficient Seismic Wave Propagation Simulation on a Low-Power Manycore Processor	57
<i>Márcio Castro, Fabrice Dupros, Emilio Franceschini, Jean-François Méhautk, and Philippe O.A. Navaux</i>	
Performance-Aware Task Management and Frequency Scaling in Embedded Systems	65
<i>Francisco Gaspar, Aleksandar Ilic, Pedro Tomás, and Leonel Sousa</i>	
Analyzing Performance Improvements and Energy Savings in Infiniband Architecture using Network Compression	73
<i>Branimir Dickov, Miquel Pericàs, Paul M. Carpenter, Nacho Navarro, and Eduard Ayguadé</i>	

Hardware Accelerators I

Bit-Parallel Approximate Pattern Matching on the Xeon Phi Coprocessor	81
<i>Tuan Tu Tran, Simon Schindel, Yongchao Liu, and Bertil Schmidt</i>	
Efficient Execution of Microscopy Image Analysis on CPU, GPU, and MIC Equipped Cluster Systems	89
<i>G. Andrade, R. Ferreira, George Teodoro, Leonardo Rocha, Joel H. Saltz, and Tahsin Kurc</i>	
High-Performance Traffic Classification on GPU	97
<i>Shijie Zhou, Prashant Rao Nittoor, and Viktor K. Prasanna</i>	
Accelerating Curvature Estimate in 3D Seismic Data Using GPGPU	105
<i>Leonardo Martins, Marco Aurélio Gonçalves da Silva, Marcelo Arruda, Joner Duarte Jr, Pedro Mário Silva, Roberto Beauclair Seixas, and Marcelo Gattass</i>	
Leveraging OmpSs to Exploit Hardware Accelerators	112
<i>Florentino Sainz, Sergi Mateo, Vicenç Beltran, Jose L. Bosque, Xavier Martorell, and Eduard Ayguadé</i>	

Transactional Memory

Improving Signature Behavior by Irrevocability in Transactional Memory Systems	120
<i>Ricardo Quisilant, Eladio Gutierrez, Emilio L. Zapata, and Oscar Plata</i>	
Scalability Analysis of Signatures in Transactional Memory Systems	128
<i>Ricardo Quisilant, Eladio Gutierrez, and Oscar Plata</i>	
Profiling Patterns of Bit Flipping for Software Transactional Memories	136
<i>Felipe L. Teixeira, Maurício L. Pilla, André R. Du Bois, and Daniel Mossé</i>	

Multi-dimensional Evaluation of Haswell's Transactional Memory Performance	144
<i>Marcio Machado Pereira, Matthew Gaudet, José Nelson Amaral, and Guido Araújo</i>	

DeTrans: Deterministic and Parallel execution of Transactions	152
<i>Vesna Smiljković, Srđan Stipić, Christof Fetzer, Osman Ünsal, Adrián Cristal, and Mateo Valero</i>	

Hardware Accelerators II

Design Space Exploration of Memory Model for Heterogeneous Computing	160
<i>Jieun Lim and Hyesoon Kim</i>	

Runtime Support for Adaptive Spatial Partitioning and Inter-Kernel Communication on GPUs	168
<i>Yash Ukidave, Charu Kalra, David Kaeli, Perhaad Mistry, and Dana Schaa</i>	

Automatic Generation of Custom Parallel Processors for Morphological Image Processing	176
<i>Emerson Carlos Pedrino and Marcio Merino Fernandes</i>	

ASIC Design of Shared Vector Accelerators for Multicore Processors	182
<i>Spiridon F. Beldianu and Sotirios G. Ziavras</i>	

Efficient Scan Operator Methods on a GPU	190
<i>Adrián P. Diéguez, Margarita Amor, and Ramón Doallo</i>	

Memory Hierarchy

Optimizing Memory Locality Using a Locality-Aware Page Table	198
<i>Eduardo H.M. Cruz, Matthias Diener, Marco A.Z. Alves, Laércio L. Pilla, and Philippe O.A. Navaux</i>	

HiTS: A High Throughput Memory Scheduling Scheme to Mitigate Denial-of-Service Attacks in Multi-core Systems	206
<i>Mansour Shafaei and Yunsi Fei</i>	

Hardware/Software Helper Thread Prefetching on Heterogeneous Many Cores	214
<i>Bharath Narasimha Swamy, Alain Ketterlin, and André Seznec</i>	

Chip-Level Optimizations

Profiling and Reducing Micro-Architecture Bottlenecks at the Hardware Level	222
<i>Francis B. Moreira, Marco A.Z. Alves, and Israel Koren</i>	

Design and Evaluation of Hierarchical Rings with Deflection Routing	230
<i>Rachata Ausavarungnirun, Chris Fallin, Xiangyao Yu, Kevin Kai-Wei Chang, Greg Nazario, Reetuparna Das, Gabriel H. Loh, and Onur Mutlu</i>	

Block Disabling Characterization and Improvements in CMPs Operating at Ultra-low Voltages	238
<i>Alexandra Ferrerón, Darío Suárez-Gracia, Jesús Alastruey-Benedé, Teresa Monreal, and Víctor Viñals</i>	

Fine-Grained Parallelism

Evaluation of a Feature Tracking Vision Application on a Heterogeneous Chip	246
<i>Rubén Gran, August Shi, Ehsan Toton, and María J. Garzarán</i>	
Predicting GPU Performance from CPU Runs Using Machine Learning	254
<i>Ioana Baldini, Stephen J. Fink, and Erik Altman</i>	
Leveraging Task-Parallelism with OmpSs in ILUPACK's Preconditioned CG Method	262
<i>J.I. Aliaga, R.M. Badía, M. Barreda, M. Bollhöfer, and E.S. Quintana-Ortí</i>	
Reducing Compiler-Inserted Instrumentation in Unified-Parallel-C Code Generation	270
<i>Michail Alvanos, José Nelson Amaral, Ettore Tiotto, Montse Farreras, and Xavier Martorell</i>	
Lock-Free GaussSieve for Linear Speedups in Parallel High Performance SVP Calculation	278
<i>Artur Mariano, Shahar Timnat, and Christian Bischof</i>	

Grid and Cloud Computing

Wide Area BonjourGrid as a Data Desktop Grid: Modeling and Implementation on Top of Redis	286
<i>Walid Saad, Leila Abidi, Heithem Abbes, Christophe Cérin, and Mohamed Jemni</i>	
An Anomaly Detection Approach for Scale-Out Storage Systems	294
<i>Guthemberg Silvestre, Carla Sauvanaud, Mohamed Kaâniche, and Karama Kanoun</i>	
Analyzing Real Cluster Data for Formulating Allocation Algorithms in Cloud Platforms	302
<i>Olivier Beaumont, Lionel Eyraud-Dubois, and Juan-Angel Lorenzo-del-Castillo</i>	
Modeling the Impact of Workload on Cloud Resource Scaling	310
<i>Anshul Gandhi, Parijat Dube, Alexei Karve, Andrzej Kochut, and Li Zhang</i>	

Parallel and Distributed Systems

SLURM Support for Remote GPU Virtualization: Implementation and Performance Study	318
<i>Sergio Iserte, Adrián Castelló, Rafael Mayo, Enrique S. Quintana-Ortí, Federico Silla, Jose Duato, Carlos Reaño, and Javier Prades</i>	
Cross-Layer Self-Adaptive/Self-Aware System Software for Exascale Systems	326
<i>Roberto Gioiosa, Gokcen Kestor, Darren J. Kerbyson, and Adolffy Hoisie</i>	

Cloud-Based OpenMP Parallelization Using a MapReduce Runtime	334
<i>Rodolfo Wottrich, Rodolfo Azevedo, and Guido Araújo</i>	
Author Index	342