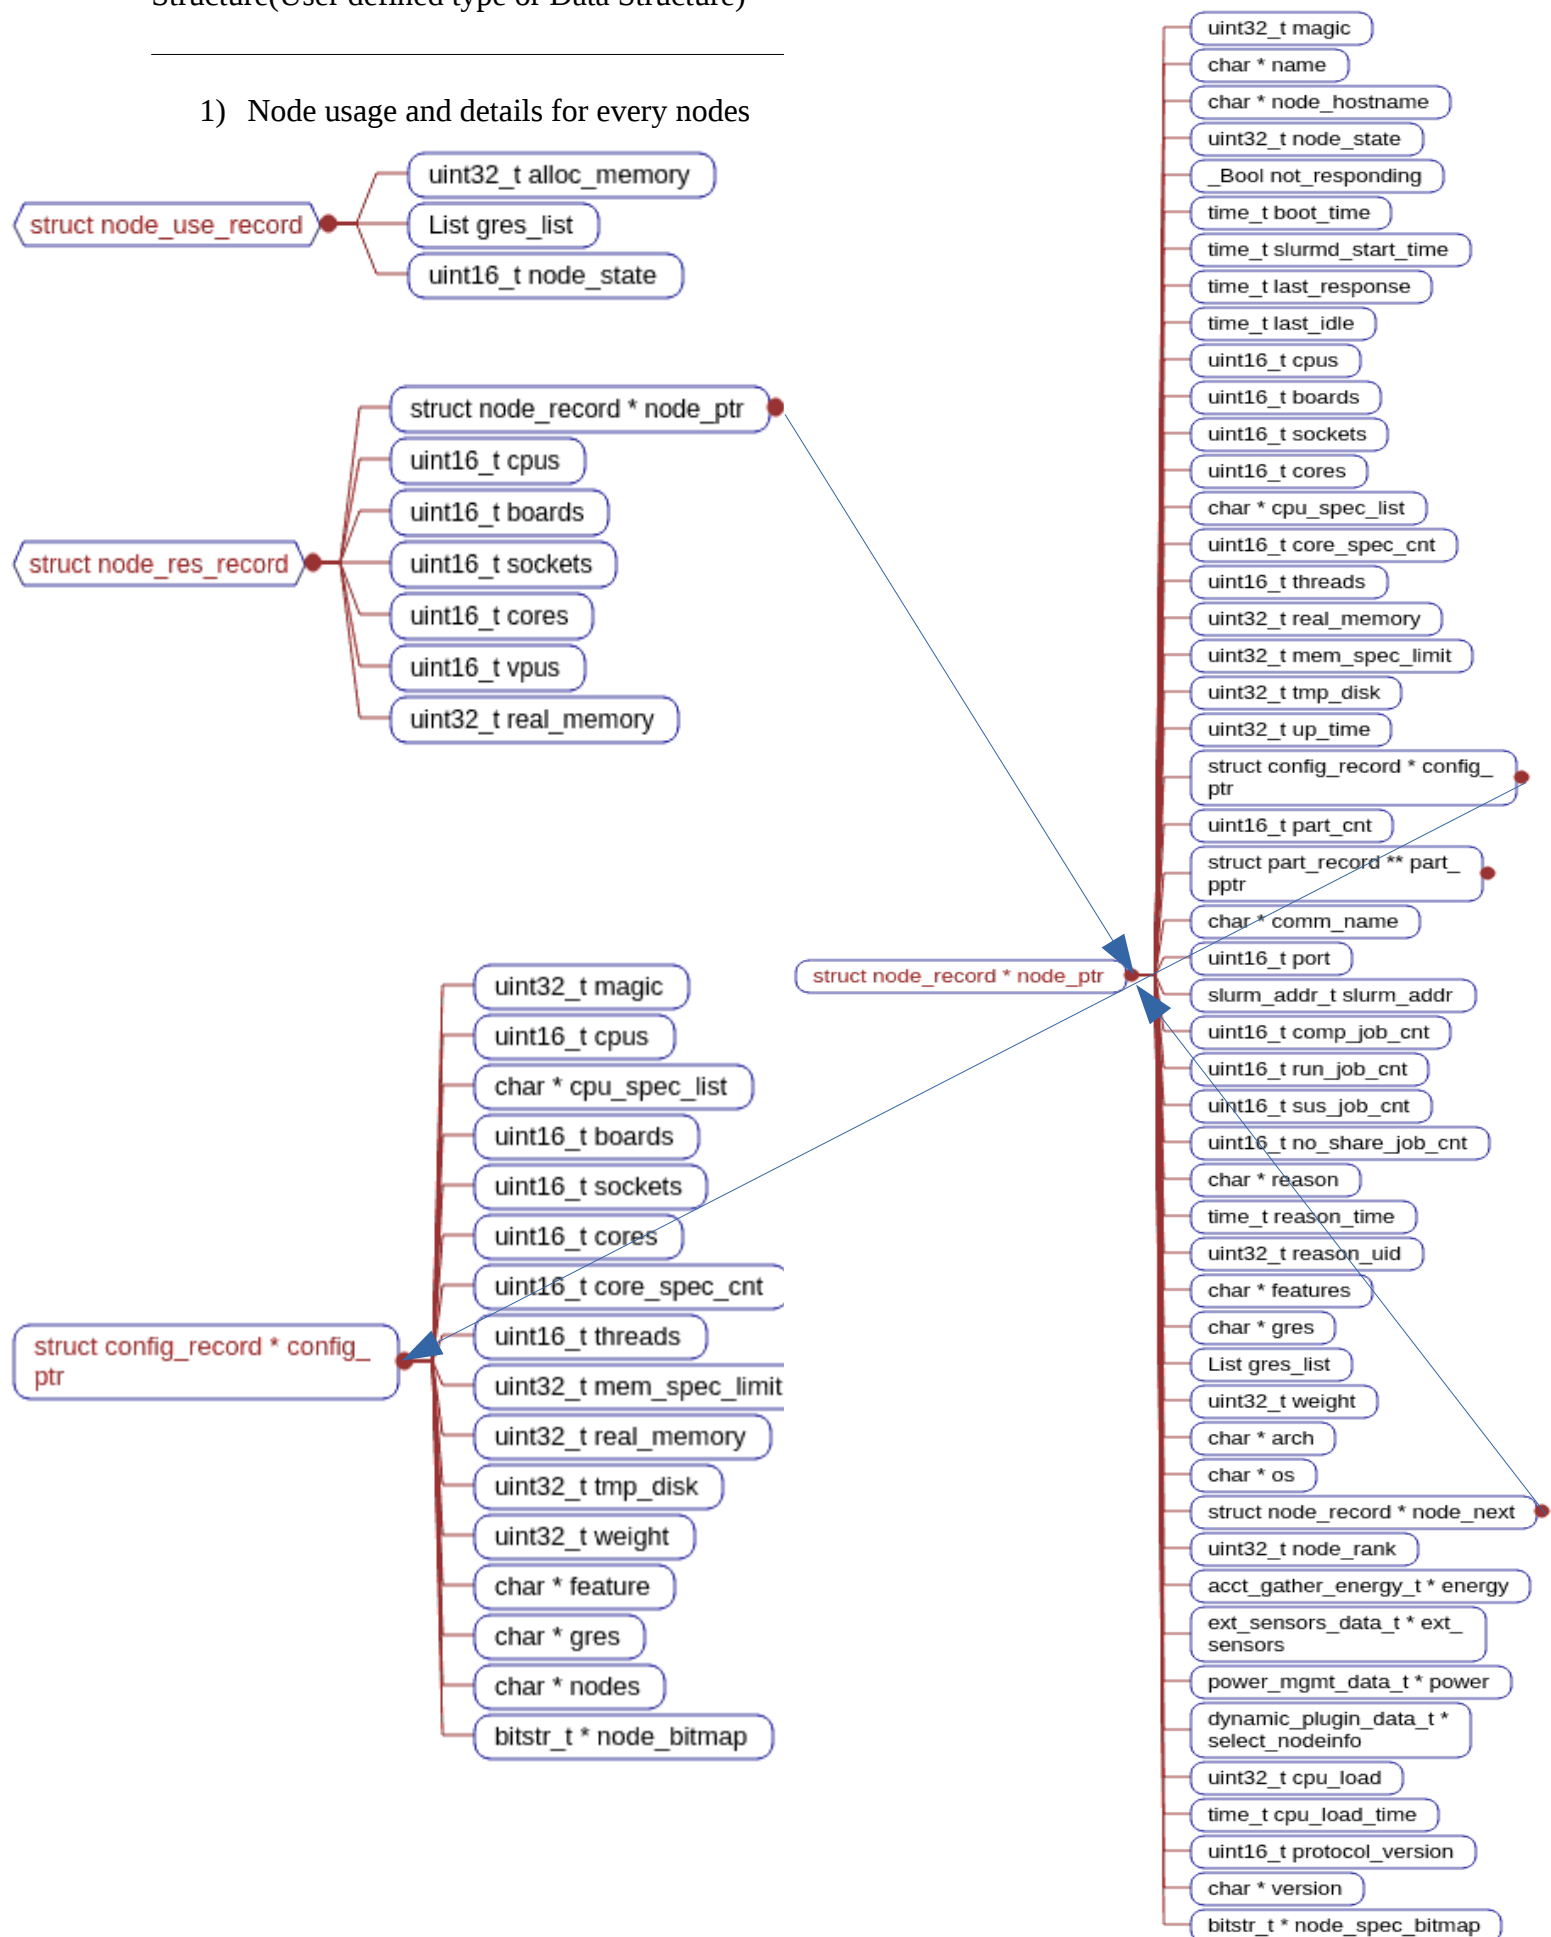
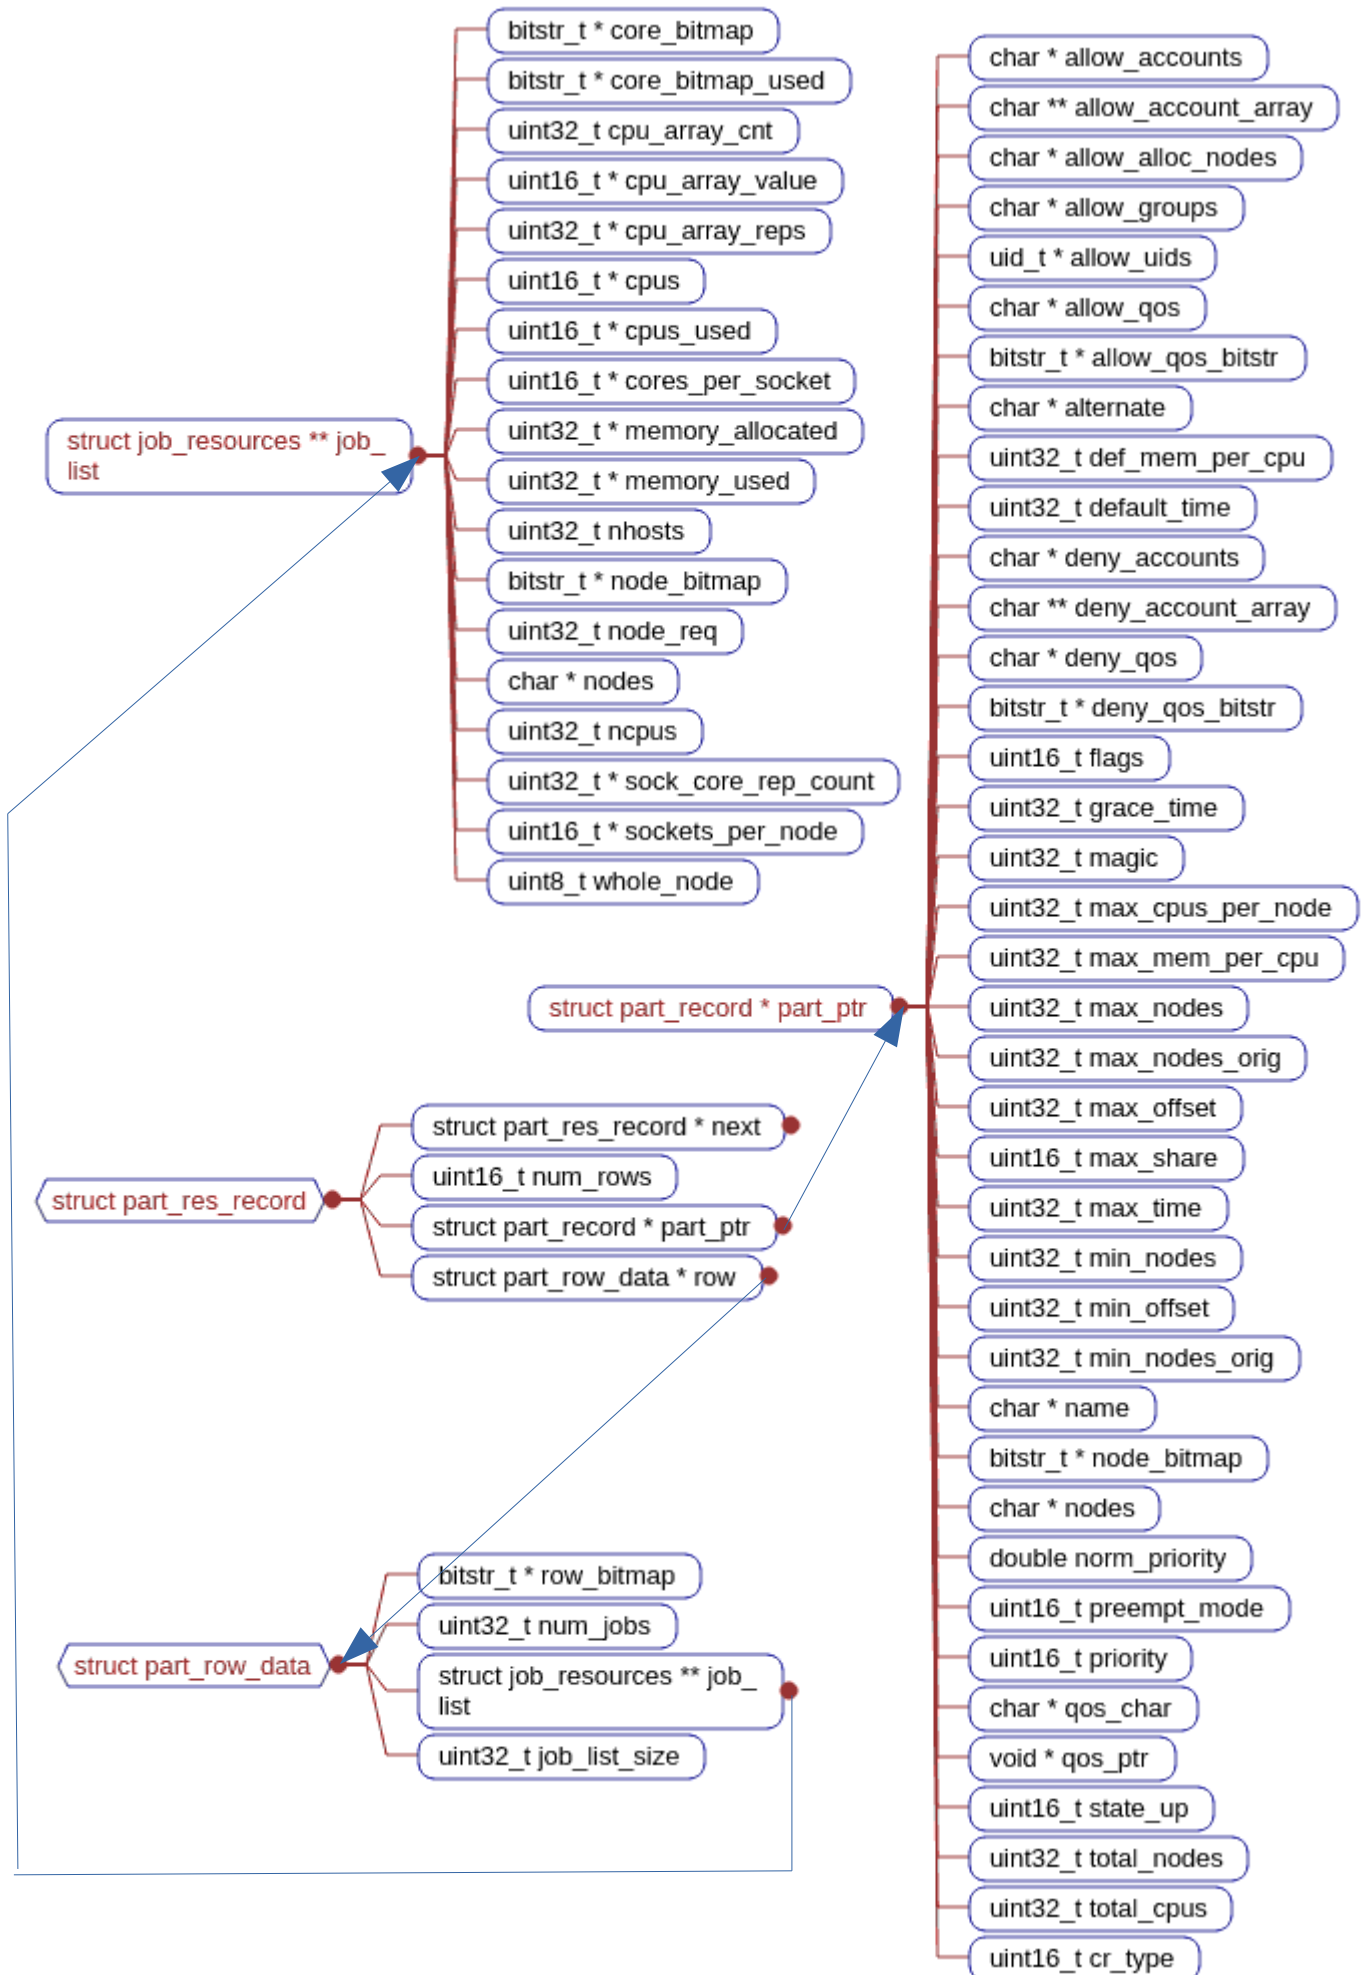


Structure(User defined type or Data Structure)

1) Node usage and details for every nodes



2) Partition and Job_resources details



```

/* struct job_resources defines exactly which resources are allocated
*    to a job, step, partition, etc.
*
* core_bitmap          - Bitmap of allocated cores for all nodes and sockets
* core_bitmap_used     - Bitmap of cores allocated to job steps
* cores_per_socket     - Count of cores per socket on this node, build by
*                        build_job_resources() and insures consistent
*                        interpretation of core_bitmap
* cpus                 - Count of desired/allocated CPUs per node for job/step
* cpus_used            - For a job, count of CPUs per node used by job steps
* cpu_array_cnt        - Count of elements in cpu_array_* below
* cpu_array_value      - Count of allocated CPUs per node for job
* cpu_array_reps       - Number of consecutive nodes on which cpu_array_value
*                        is duplicated. See NOTES below.
* memory_allocated     - MB per node reserved for the job or step
* memory_used          - MB per node of memory consumed by job steps
* nhosts               - Number of nodes in the allocation. On a
*                        bluegene machine this represents the number
*                        of midplanes used. This should always be
*                        the number of bits set in node_bitmap.
* node_bitmap          - Bitmap of nodes allocated to the job. Unlike the
*                        node_bitmap in slurmctld's job record, the bits
*                        here do NOT get cleared as the job completes on a
*                        node
* node_req             - NODE_CR_RESERVED|NODE_CR_ONE_ROW|
NODE_CR_AVAILABLE
* nodes               - Names of nodes in original job allocation
* ncpus               - Number of processors in the allocation
* sock_core_rep_count  - How many consecutive nodes that sockets_per_node
*                        and cores_per_socket apply to, build by
*                        build_job_resources() and insures consistent
*                        interpretation of core_bitmap
* sockets_per_node     - Count of sockets on this node, build by
*                        build_job_resources() and insures consistent
*                        interpretation of core_bitmap
* whole_node          - Job allocated full node (used only by select/cons_res)
*
* NOTES:
* cpu_array_* contains the same information as "cpus", but in a more compact
* format. For example if cpus = {4, 4, 2, 2, 2, 2, 2, 2} then cpu_array_cnt=2
* cpu_array_value = {4, 2} and cpu_array_reps = {2, 6}. We do not need to
* save/restore these values, but generate them by calling
* build_job_resources_cpu_array()
*
* Sample layout of core_bitmap:
* |      Node_0      |      Node_1      |
* |  Sock_0  |  Sock_1  |  Sock_0  |  Sock_1  |
* | Core_0 | Core_1 | Core_0 | Core_1 | Core_0 | Core_1 | Core_0 | Core_1 |
* | Bit_0  | Bit_1  | Bit_2  | Bit_3  | Bit_4  | Bit_5  | Bit_6  | Bit_7  |
*
* If a job changes size (relinquishes nodes), the node_bitmap will remain
* unchanged, but cpus, cpus_used, cpus_array_*, and memory_used will be

```

* updated (e.g. cpus and mem_used on that node cleared).

*/

struct job_resources

Node-0				Node-1			
Sock-0		Sock-1		Sock-2		Sock-3	
Core-0	Core-1	Core-2	Core-3	Core-4	Core-5	Core-6	Core-7

/*

* node_res_record.node_state assists with the unique state of each node.

* When a job is allocated, these flags provide protection for nodes in a

* Shared=NO or Shared=EXCLUSIVE partition from other jobs.

*

* NOTES:

* - If node is in use by Shared=NO part, some CPUs/memory may be available

* - Caution with NODE_CR_AVAILABLE: a Sharing partition could be full.

*

* - these values are staggered so that they can be incremented as multiple

* jobs are allocated to each node. This is needed to be able to support

* preemption, which can override these protections.

*/

enum node_cr_state {

NODE_CR_AVAILABLE = 0, /* The node may be IDLE or IN USE (shared) */

NODE_CR_ONE_ROW = 1, /* node is in use by Shared=NO part */

NODE_CR_RESERVED = 64000 /* node is in use by Shared=EXCLUSIVE part */

};

/* job_details - specification of a job's constraints,

* can be purged after initiation */

struct job_details {

char *acctg_freq; /* accounting polling interval */

uint32_t argc; /* count of argv elements */

char **argv; /* arguments for a batch job script */

time_t begin_time; /* start at this time (srun --begin),

* resets to time first eligible

* (all dependencies satisfied) */

char *ckpt_dir; /* directory to store checkpoint

* images */

uint16_t contiguous; /* set if requires contiguous nodes */

uint16_t core_spec; /* specialized core count */

char *cpu_bind; /* binding map for map/mask_cpu - This

* currently does not matter to the

* job allocation, setting this does

* not do anything for steps. */

uint16_t cpu_bind_type; /* see cpu_bind_type_t - This

* currently does not matter to the

* job allocation, setting this does

```

uint32_t cpu_freq_min; /* not do anything for steps. */
uint32_t cpu_freq_max; /* Minimum cpu frequency */
uint32_t cpu_freq_gov; /* Maximum cpu frequency */
uint16_t cpus_per_task; /* cpu frequency governor */
                        /* number of processors required for
                        * each task */
List depend_list; /* list of job_ptr:state pairs */
char *dependency; /* wait for other jobs */
char *orig_dependency; /* original value (for archiving) */
uint16_t env_cnt; /* size of env_sup (see below) */
char **env_sup; /* supplemental environment variables
                * as set by Moab */
bitstr_t *exc_node_bitmap; /* bitmap of excluded nodes */
char *exc_nodes; /* excluded nodes */
uint32_t expanding_jobid; /* ID of job to be expanded */
List feature_list; /* required features with
                  * node counts */
char *features; /* required features */
uint32_t magic; /* magic cookie for data integrity */
uint32_t max_cpus; /* maximum number of cpus */
uint32_t max_nodes; /* maximum number of nodes */
multi_core_data_t *mc_ptr; /* multi-core specific data */
char *mem_bind; /* binding map for map/mask_cpu */
uint16_t mem_bind_type; /* see mem_bind_type_t */
uint32_t min_cpus; /* minimum number of cpus */
uint32_t min_nodes; /* minimum number of nodes */
uint16_t nice; /* requested priority change,
               * NICE_OFFSET == no change */
uint16_t ntasks_per_node; /* number of tasks on each node */
uint32_t num_tasks; /* number of tasks to start */
uint8_t open_mode; /* stdout/err append or truncate */
uint8_t overcommit; /* processors being over subscribed */
uint16_t plane_size; /* plane size when task_dist =
                     * SLURM_DIST_PLANE */

/* job constraints: */
uint32_t pn_min_cpus; /* minimum processors per node */
uint32_t pn_min_memory; /* minimum memory per node (MB) OR
                        * memory per allocated
                        * CPU | MEM_PER_CPU */
uint32_t pn_min_tmp_disk; /* minimum tempdisk per node, MB */
uint8_t prolog_running; /* set while prolog_slurmctld is
                        * running */
uint32_t reserved_resources; /* CPU minutes of resources reserved
                             * for this job while it was pending */
bitstr_t *req_node_bitmap; /* bitmap of required nodes */
uint16_t *req_node_layout; /* task layout for required nodes */
time_t preempt_start_time; /* time that preemption began to start
                           * this job */
char *req_nodes; /* required nodes */
uint16_t requeue; /* controls ability requeue job */
char *restart_dir; /* restart execution from ckpt images
                  * in this dir */

```

```

uint8_t share_res;          /* set if job can share resources with
                             * other jobs */
char *std_err;              /* pathname of job's stderr file */
char *std_in;               /* pathname of job's stdin file */
char *std_out;              /* pathname of job's stdout file */
time_t submit_time;         /* time of submission */
uint16_t task_dist;         /* task layout for this job. Only
                             * useful when Consumable Resources
                             * is enabled */

uint32_t usable_nodes;      /* node count needed by preemption */
uint8_t whole_node;         /* job requested exclusive node use */
char *work_dir;             /* pathname of working directory */
};

```

```

typedef struct job_array_struct {
    uint32_t task_cnt;        /* count of remaining task IDs */
    bitstr_t *task_id_bitmap; /* bitmap of remaining task IDs */
    char *task_id_str;        /* string describing remaining task IDs,
                             * needs to be recalculated if NULL */

    uint32_t array_flags;     /* Flags to control behavior (FUTURE) */
    uint32_t max_run_tasks;    /* Maximum number of running tasks */
    uint32_t tot_run_tasks;    /* Current running task count */
    uint32_t min_exit_code;    /* Minimum exit code from any task */
    uint32_t max_exit_code;    /* Maximum exit code from any task */
    uint32_t tot_comp_tasks;   /* Completed task count */
} job_array_struct_t;

```

```

/*
 * NOTE: When adding fields to the job_record, or any underlying structures,
 * be sure to sync with _rec_job_copy.
 */

```

```

struct job_record {
    char *account;            /* account number to charge */
    char *alias_list;         /* node name to address aliases */
    char *alloc_node;         /* local node making resource alloc */
    uint16_t alloc_resp_port; /* RESPONSE_RESOURCE_ALLOCATION port */
    uint32_t alloc_sid;       /* local sid making resource alloc */
    uint32_t array_job_id;     /* job_id of a job array or 0 if N/A */
    uint32_t array_task_id;    /* task_id of a job array */
    job_array_struct_t *array_recs; /* job array details,
                                     * only in meta-job record */

    uint32_t assoc_id;         /* used for accounting plugins */
    void *assoc_ptr;           /* job's assoc record ptr, it is
                             * void* because of interdependencies
                             * in the header files, confirm the
                             * value before use */

    uint16_t batch_flag;       /* 1 or 2 if batch job (with script),
                             * 2 indicates retry mode (one retry) */

    char *batch_host;          /* host executing batch script */
    char *burst_buffer;        /* burst buffer specification */
    check_jobinfo_t check_job; /* checkpoint context, opaque */
    uint16_t ckpt_interval;    /* checkpoint interval in minutes */
}

```

```

time_t ckpt_time;          /* last time job was periodically
                           * checkpointed */

char *comment;             /* arbitrary comment */

uint32_t cpu_cnt;          /* current count of CPUs held
                           * by the job, decremented while job is
                           * completing (N/A for bluegene
                           * systems) */

uint16_t cr_enabled;       /* specify if Consumable Resources
                           * is enabled. Needed since CR deals
                           * with a finer granularity in its
                           * node/cpu scheduling (available cpus
                           * instead of available nodes) than the
                           * bluegene and the linear plugins
                           * 0 if cr is NOT enabled,
                           * 1 if cr is enabled */

uint32_t db_index;         /* used only for database
                           * plugins */

uint32_t derived_ec;       /* highest exit code of all job steps */

struct job_details *details; /* job details */

uint16_t direct_set_prio;   /* Priority set directly if
                           * set the system will not
                           * change the priority any further. */

time_t end_time;           /* time execution ended, actual or
                           * expected. if terminated from suspend
                           * state, this is time suspend began */

bool epilog_running;       /* true if EpilogSlurmctld is running */

uint32_t exit_code;        /* exit code for job (status from
                           * wait call) */

front_end_record_t *front_end_ptr; /* Pointer to front-end node running
                                   * this job */

char *gres;                /* generic resources requested by job */
List gres_list;            /* generic resource allocation detail */
char *gres_alloc;          /* Allocated GRES added over all nodes
                           * to be passed to slurmdbd */

char *gres_req;            /* Requested GRES added over all nodes
                           * to be passed to slurmdbd */

char *gres_used;           /* Actual GRES use added over all nodes
                           * to be passed to slurmdbd */

uint32_t group_id;         /* group submitted under */

uint32_t job_id;           /* job ID */

struct job_record *job_next; /* next entry with same hash index */
struct job_record *job_array_next_j; /* job array linked list by job_id */
struct job_record *job_array_next_t; /* job array linked list by task_id */
job_resources_t *job_resrcs; /* details of allocated cores */

uint16_t job_state;        /* state of the job */

uint16_t kill_on_node_fail; /* 1 if job should be killed on
                           * node failure */

char *licenses;            /* licenses required by the job */
List license_list;         /* structure with license info */

uint16_t limit_set_max_cpus; /* if max_cpus was set from
                           * a limit false if user set */

uint16_t limit_set_max_nodes; /* if max_nodes was set from

```

```

                                /* a limit false if user set */
uint16_t limit_set_min_cpus; /* if max_cpus was set from
                                /* a limit false if user set */
uint16_t limit_set_min_nodes; /* if max_nodes was set from
                                /* a limit false if user set */
uint16_t limit_set_pn_min_memory; /* if pn_min_memory was set from
                                /* a limit false if user set */
uint16_t limit_set_time; /* if time_limit was set from
                                /* a limit false if user set */
uint16_t limit_set_qos; /* if qos_limit was set from
                                /* a limit false if user set */
uint16_t mail_type; /* see MAIL_JOB_ in slurm.h */
char *mail_user; /* user to get e-mail notification */
uint32_t magic; /* magic cookie for data integrity */
char *name; /* name of the job */
char *network; /* network/switch requirement spec */
uint32_t next_step_id; /* next step id to be used */
char *nodes; /* list of nodes allocated to job */
slurm_addr_t *node_addr; /* addresses of the nodes allocated to
                                * job */
bitstr_t *node_bitmap; /* bitmap of nodes allocated to job */
bitstr_t *node_bitmap_cg; /* bitmap of nodes completing job */
uint32_t node_cnt; /* count of nodes currently
                                * allocated to job */
uint32_t node_cnt_wag; /* count of nodes Slurm thinks
                                * will be allocated when the
                                * job is pending and node_cnt
                                * wasn't given by the user.
                                * This is packed in total_nodes
                                * when dumping state. When
                                * state is read in check for
                                * pending state and set this
                                * instead of total_nodes */
char *nodes_completing; /* nodes still in completing state
                                * for this job, used to insure
                                * epilog is not re-run for job */
uint16_t other_port; /* port for client communications */
char *partition; /* name of job partition(s) */
List part_ptr_list; /* list of pointers to partition recs */
bool part_nodes_missing; /* set if job's nodes removed from this
                                * partition */
struct part_record *part_ptr; /* pointer to the partition record */
uint8_t power_flags; /* power management flags,
                                * see SLURM_POWER_FLAGS */
time_t pre_sus_time; /* time job ran prior to last suspend */
time_t preempt_time; /* job preemption signal time */
bool preempt_in_progress; /* Preemption of other jobs in progress
                                * in order to start this job,
                                * (Internal use only, don't save) */
uint32_t priority; /* relative priority of the job,
                                * zero == held (don't initiate) */
uint32_t *priority_array; /* partition based priority */

```



```

priority_factors_object_t *prio_factors; /* cached value used
                                         * by sprio command */
uint32_t profile; /* Acct_gather_profile option */
uint32_t qos_id; /* quality of service id */
void *qos_ptr; /* pointer to the quality of
               * service record used for
               * this job, it is
               * void* because of interdependencies
               * in the header files, confirm the
               * value before use */

uint8_t reboot; /* node reboot requested before start */
uint16_t restart_cnt; /* count of restarts */
time_t resize_time; /* time of latest size change */
uint32_t resv_id; /* reservation ID */
char *resv_name; /* reservation name */
struct slurmctld_resv *resv_ptr; /* reservation structure pointer */
uint32_t requid; /* requester user ID */
char *resp_host; /* host for srun communications */
char *sched_nodes; /* list of nodes scheduled for job */
dynamic_plugin_data_t *select_jobinfo; /* opaque data, BlueGene */
uint8_t sicp_mode; /* set for inter-cluster jobs */
char **spank_job_env; /* environment variables for job prolog
                      * and epilog scripts as set by SPANK
                      * plugins */

uint32_t spank_job_env_size; /* element count in spank_env */
uint16_t start_protocol_ver; /* Slurm version step was
                             * started with */

time_t start_time; /* time execution begins,
                  * actual or expected */

char *state_desc; /* optional details for state_reason */
uint32_t state_reason; /* reason job still pending or failed
                       * see slurm.h:enum job_wait_reason */

List step_list; /* list of job's steps */
time_t suspend_time; /* time job last suspended or resumed */
time_t time_last_active; /* time of last job activity */
uint32_t time_limit; /* time_limit minutes or INFINITE,
                    * NO_VAL implies partition max_time */

uint32_t time_min; /* minimum time_limit minutes or
                  * INFINITE,
                  * zero implies same as time_limit */

time_t tot_sus_time; /* total time in suspend state */
uint32_t total_cpus; /* number of allocated cpus,
                    * for accounting */

uint32_t total_nodes; /* number of allocated nodes
                     * for accounting */

uint32_t user_id; /* user the job runs as */
uint16_t wait_all_nodes; /* if set, wait for all nodes to boot
                         * before starting the job */

uint16_t warn_flags; /* flags for signal to send */
uint16_t warn_signal; /* signal to send before end_time */
uint16_t warn_time; /* when to send signal before
                   * end_time (secs) */

```

```

char *wckey; /* optional wckey */

/* Request number of switches support */
uint32_t req_switch; /* Minimum number of switches */
uint32_t wait4switch; /* Maximum time to wait for minimum switches */
bool best_switch; /* true=min number of switches met */
time_t wait4switch_start; /* Time started waiting for switch */
};

/*****\
* PARTITION parameters and data structures
\*****/
#define PART_MAGIC 0xae8495

struct part_record {
    char *allow_accounts; /* comma delimited list of accounts,
                          * NULL indicates all */
    char **allow_account_array; /* NULL terminated list of allowed
                              * accounts */
    char *allow_alloc_nodes; /* comma delimited list of allowed
                              * allocating nodes
                              * NULL indicates all */
    char *allow_groups; /* comma delimited list of groups,
                       * NULL indicates all */
    uid_t *allow_uids; /* zero terminated list of allowed user IDs */
    char *allow_qos; /* comma delimited list of qos,
                    * NULL indicates all */
    bitstr_t *allow_qos_bitstr; /* (DON'T PACK) associated with
                              * char *allow_qos but used internally */
    char *alternate; /* name of alternate partition */
    uint32_t def_mem_per_cpu; /* default MB memory per allocated CPU */
    uint32_t default_time; /* minutes, NO_VAL or INFINITE */
    char *deny_accounts; /* comma delimited list of denied accounts */
    char **deny_account_array; /* NULL terminated list of denied accounts */
    char *deny_qos; /* comma delimited list of denied qos */
    bitstr_t *deny_qos_bitstr; /* (DON'T PACK) associated with
                              * char *deny_qos but used internally */
    uint16_t flags; /* see PART_FLAG_* in slurm.h */
    uint32_t grace_time; /* default preempt grace time in seconds */
    uint32_t magic; /* magic cookie to test data integrity */
    uint32_t max_cpus_per_node; /* maximum allocated CPUs per node */
    uint32_t max_mem_per_cpu; /* maximum MB memory per allocated CPU */
    uint32_t max_nodes; /* per job or INFINITE */
    uint32_t max_nodes_orig; /* unscaled value (c-nodes on BlueGene) */
    uint32_t max_offset; /* select plugin max offset */
    uint16_t max_share; /* number of jobs to gang schedule */
    uint32_t max_time; /* minutes or INFINITE */
    uint32_t min_nodes; /* per job */
    uint32_t min_offset; /* select plugin min offset */
    uint32_t min_nodes_orig; /* unscaled value (c-nodes on BlueGene) */
    char *name; /* name of the partition */
    bitstr_t *node_bitmap; /* bitmap of nodes in partition */

```

```

char *nodes;          /* comma delimited list names of nodes */
double norm_priority; /* normalized scheduling priority for
                      * jobs (DON'T PACK) */
uint16_t preempt_mode; /* See PREEMPT_MODE_* in slurm/slurm.h */
uint16_t priority;     /* scheduling priority for jobs */
char *qos_char;        /* requested QOS from slurm.conf */
void *qos_ptr;         /* pointer to the quality of
                      * service record attached to this
                      * partition, it is void* because of
                      * interdependencies in the header
                      * files, confirm the value before use */
uint16_t state_up;     /* See PARTITION_* states in slurm.h */
uint32_t total_nodes;  /* total number of nodes in the partition */
uint32_t total_cpus;   /* total number of cpus in the partition */
uint32_t max_cpu_cnt;  /* max # of cpus on a node in the partition */
uint32_t max_core_cnt; /* max # of cores on a node in the partition */
uint16_t cr_type;      /* Custom CR values for partition (if supported by select plugin) */
};

```

```

extern List part_list;          /* list of part_record entries */
extern time_t last_part_update; /* time of last part_list update */
extern struct part_record default_part; /* default configuration values */
extern char *default_part_name; /* name of default partition */
extern struct part_record *default_part_loc; /* default partition ptr */
extern uint16_t part_max_priority; /* max priority in all partitions */

```

```

struct config_record {
    uint32_t magic;          /* magic cookie to test data integrity */
    uint16_t cpus;           /* count of processors running on the node */
    char *cpu_spec_list;     /* arbitrary list of specialized cpus */
    uint16_t boards;        /* count of boards configured */
    uint16_t sockets;       /* number of sockets per node */
    uint16_t cores;         /* number of cores per CPU */
    uint16_t core_spec_cnt;  /* number of specialized cores */
    uint16_t threads;       /* number of threads per core */
    uint32_t mem_spec_limit; /* MB real memory for memory specialization */
    uint32_t real_memory;    /* MB real memory on the node */
    uint32_t tmp_disk;       /* MB total storage in TMP_FS file system */
    uint32_t weight;         /* arbitrary priority of node for
                          * scheduling work on */

    char *feature;           /* arbitrary list of node's features */
    char *gres;              /* arbitrary list of node's generic resources */
    char *nodes;             /* name of nodes with this configuration */
    bitstr_t *node_bitmap;   /* bitmap of nodes with this configuration */
};

```

```

extern List config_list; /* list of config_record entries */

```

```

extern List front_end_list; /* list of slurm_conf_frontend_t entries */

```

```
struct node_record {
    uint32_t magic; /* magic cookie for data integrity */
    char *name; /* name of the node. NULL==defunct */
    char *node_hostname; /* hostname of the node */
    uint32_t node_state; /* enum node_states, ORed with
        * NODE_STATE_NO_RESPOND if not
        * responding */
    bool not_responding; /* set if fails to respond,
        * clear after logging this */
    time_t boot_time; /* Time of node boot,
        * computed from up_time */
    time_t slurmd_start_time; /* Time of slurmd startup */
    time_t last_response; /* last response from the node */
    time_t last_idle; /* time node last become idle */
    uint16_t cpus; /* count of processors on the node */
    uint16_t boards; /* count of boards configured */
    uint16_t sockets; /* number of sockets per node */
    uint16_t cores; /* number of cores per CPU */
    char *cpu_spec_list; /* node's specialized cpus */
    uint16_t core_spec_cnt; /* number of specialized cores on node*/
    uint16_t threads; /* number of threads per core */
    uint32_t real_memory; /* MB real memory on the node */
    uint32_t mem_spec_limit; /* MB memory limit for specialization */
    uint32_t tmp_disk; /* MB total disk in TMP_FS */
    uint32_t up_time; /* seconds since node boot */
    struct config_record *config_ptr; /* configuration spec ptr */
    uint16_t part_cnt; /* number of associated partitions */
    struct part_record **part_pptr; /* array of pointers to partitions
        * associated with this node*/
    char *comm_name; /* communications path name to node */
    uint16_t port; /* TCP port number of the slurmd */
    slurm_addr_t slurm_addr; /* network address */
    uint16_t comp_job_cnt; /* count of jobs completing on node */
    uint16_t run_job_cnt; /* count of jobs running on node */
    uint16_t sus_job_cnt; /* count of jobs suspended on node */
    uint16_t no_share_job_cnt; /* count of jobs running that will
        * not share nodes */
    char *reason; /* why a node is DOWN or DRAINING */
    time_t reason_time; /* Time stamp when reason was
        * set, ignore if no reason is set. */
    uint32_t reason_uid; /* User that set the reason, ignore if
        * no reason is set. */
    char *features; /* node's features, used only
        * for state save/restore, DO NOT
        * use for scheduling purposes */
    char *gres; /* node's generic resources, used only
        * for state save/restore, DO NOT
        * use for scheduling purposes */
    List gres_list; /* list of gres state info managed by
        * plugins */
    uint32_t weight; /* original weight, used only for state
        * save/restore, DO NOT use for
```

```

        * scheduling purposes. */
char *arch;          /* computer architecture */
char *os;            /* operating system now running */
struct node_record *node_next; /* next entry with same hash index */
uint32_t node_rank;  /* Hilbert number based on node name,
                     * or other sequence number used to
                     * order nodes by location,
                     * no need to save/restore */

#ifdef HAVE_ALPS_CRAY
    uint32_t basil_node_id; /* Cray-XT BASIL node ID,
                             * no need to save/restore */
    time_t down_time; /* When first set to DOWN state */
#endif /* HAVE_ALPS_CRAY */
    acct_gather_energy_t *energy; /* power consumption data */
    ext_sensors_data_t *ext_sensors; /* external sensor data */
    power_mgmt_data_t *power; /* power management data */
    dynamic_plugin_data_t *select_nodeinfo; /* opaque data structure,
                                             * use select_g_get_nodeinfo()
                                             * to access contents */

    uint32_t cpu_load; /* CPU load * 100 */
    time_t cpu_load_time; /* Time when cpu_load last set */
    uint16_t protocol_version; /* Slurm version number */
    char *version; /* Slurm version */
    bitstr_t *node_spec_bitmap; /* node cpu specialization bitmap */
};

extern struct node_record *node_record_table_ptr; /* ptr to node records */
extern int node_record_count; /* count in node_record_table_ptr */
extern xhash_t *node_hash_table; /* hash table for node records */
extern time_t last_node_update; /* time of last node record update */

extern uint16_t *cr_node_num_cores;
extern uint32_t *cr_node_cores_offset;

/*
 * bitmap2node_name_sortable - given a bitmap, build a list of comma
 * separated node names. names may include regular expressions
 * (e.g. "lx[01-10]")
 * IN bitmap - bitmap pointer
 * IN sort - returned ordered list or not
 * RET pointer to node list or NULL on error
 * globals: node_record_table_ptr - pointer to node table
 * NOTE: the caller must xfree the memory at node_list when no longer required
 */
char * bitmap2node_name_sortable (bitstr_t *bitmap, bool sort);

/*
 * bitmap2node_name - given a bitmap, build a list of comma separated node
 * names. names may include regular expressions (e.g. "lx[01-10]")
 * IN bitmap - bitmap pointer
 * RET pointer to node list or NULL on error
 * globals: node_record_table_ptr - pointer to node table
 * NOTE: the caller must xfree the memory at node_list when no longer required

```

```

*/
char * bitmap2node_name (bitstr_t *bitmap);

/*****
* Consumable Resources parameters and data structures
*****/

/*
* Define the type of update and of data retrieval that can happen
* from the "select/cons_res" plugin. This information needed to
* support processors as consumable resources. This structure will be
* useful when updating other types of consumable resources as well
*/
enum select_plugin_data_info {
    SELECT_CR_PLUGIN, /* data-> uint32 1 if CR plugin */
    SELECT_BITMAP, /* Unused since version 2.0 */
    SELECT_ALLOC_CPUS, /* data-> uint16 alloc cpus (CR support) */
    SELECT_ALLOC_LPS, /* data-> uint32 alloc lps (CR support) */
    SELECT_AVAIL_MEMORY, /* data-> uint32 avail mem (CR support) */
    SELECT_STATIC_PART, /* data-> uint16, 1 if static partitioning
                        * BlueGene support */
    SELECT_CONFIG_INFO /* data-> List get .conf info from select
                        * plugin */
};

job_resource.h,node_conf.h,slurmctld.h,

```